

# Database design

Part 1

# Chapter 7: Entity-Relationship Model

# Chapter 7: Entity-Relationship Model

- Design Process
- Modeling
- Mapping Constraints
- E-R Diagram
- Design Issues
- Weak Entity Sets
- Extended E-R Features
- Design of the Bank Database
- Reduction to Relation Schemas
- Database Design
- UML



# DESIGN PROCESS

PPD

- Design Process
- E-R Model



## Design Phases

- The initial phase of database design is to characterize fully the data needs of the prospective database users
- Next, the designer chooses a data model and, by applying the concepts of the chosen data model, translates these requirements into a conceptual schema of the database
- A fully developed conceptual schema also indicates the functional requirements of the enterprise. In a "specification of functional requirements", users describe the kinds of operations (or transactions) that will be performed on the data



## Design Phases (Cont.)

The process of moving from an abstract data model to the implementation of the database proceeds in two final design phases.

- Logical Design – Deciding on the database schema.  
Database design requires that we find a “good” collection of relation schemas.
  - Business decision – What attributes should we record in the database?
  - Computer Science decision – What relation schemas should we have and how should the attributes be distributed among the various relation schemas?
- Physical Design – Deciding on the physical layout of the database



## Design Approaches

- Entity Relationship Model (covered in this chapter)
  - Models an enterprise as a collection of *entities* and *relationships*
    - Entity: a “thing” or “object” in the enterprise that is distinguishable from other objects
      - Described by a set of *attributes*
    - Relationship: an association among several entities
      - Represented diagrammatically by an *entity-relationship diagram*:
  - Normalization Theory (Chapter 8)
    - Formalize what designs are bad, and test for them



PPD

- Design Process
- E-R Model

## E-R MODEL

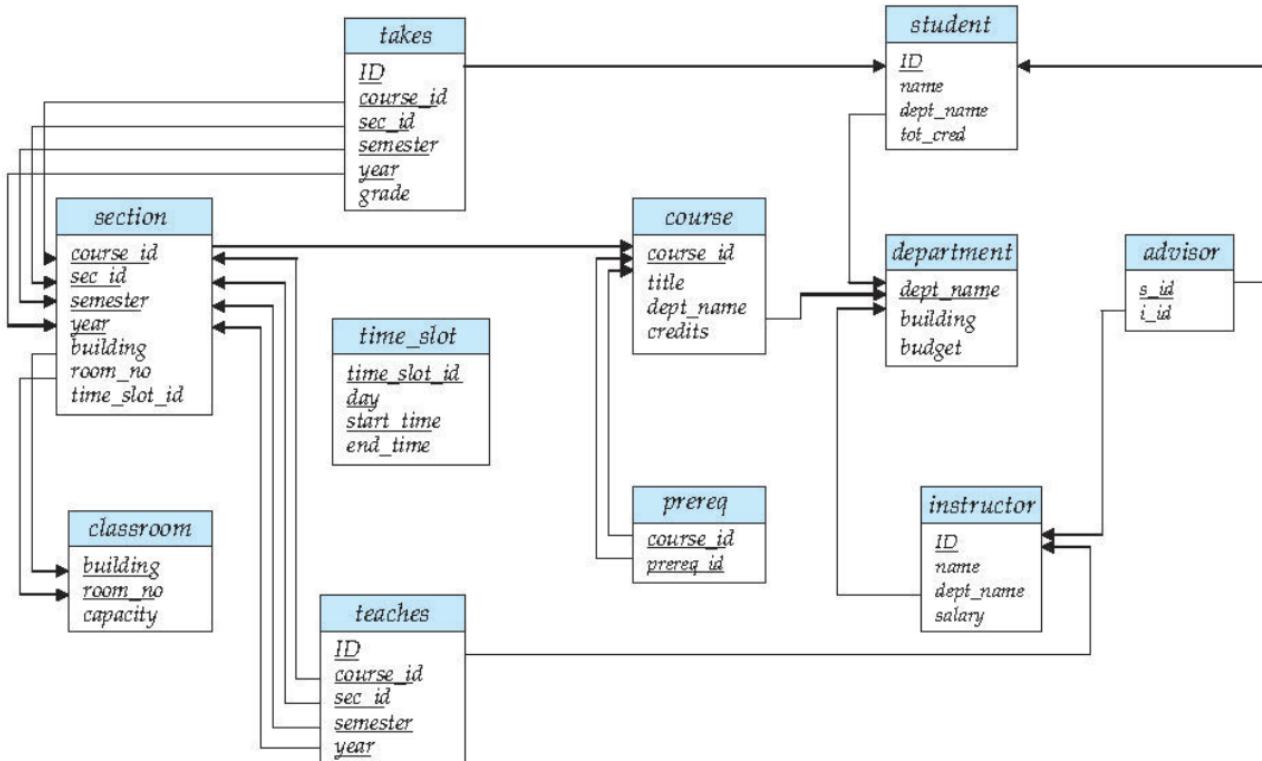


## ER model – Database Modeling

- The ER data model was developed to facilitate database design by allowing specification of an **enterprise schema** that represents the overall logical structure of a database
- The ER model is very useful in mapping the meanings and interactions of real-world enterprises onto a conceptual schema. Because of this usefulness, many database-design tools draw on concepts from the ER model
- The ER data model employs three basic concepts:
  - entity sets
  - relationship sets
  - attributes
- The ER model also has an associated diagrammatic representation, the ER diagram, which can express the overall logical structure of a database graphically

# Modeling

- A *database* can be modeled as:
  - a collection of entities,
  - relationship among entities.
- An **entity** is an object that exists and is distinguishable from other objects.
  - Example: person, company, event, plant, Course, Department
- Entities have **attributes**
  - Example: people have *names* and *addresses*
- An **entity set** is a **set of entities of the same type** that share the same properties.
  - Example: set of all persons, companies, trees, holidays



# Entity Sets *instructor* and *student*

*set of tuples*

instructor_ID	instructor_name
76766	Crick
45565	Katz
10101	Srinivasan
98345	Kim
76543	Singh
22222	Einstein

*instructor*

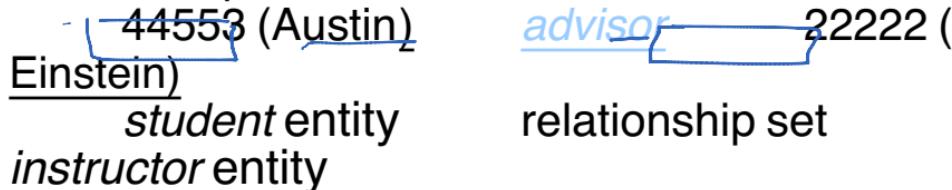
student-ID	student_name
98988	Tanaka
12345	Shankar
00128	Zhang
76543	Brown
76653	Aoi
23121	Chavez
44553	Peltier

*student*

# Relationship Sets

- A **relationship** is an **association among several entities**

Example:



- A **relationship set** is a mathematical relation among  $n \geq 2$  entities, each taken from entity sets

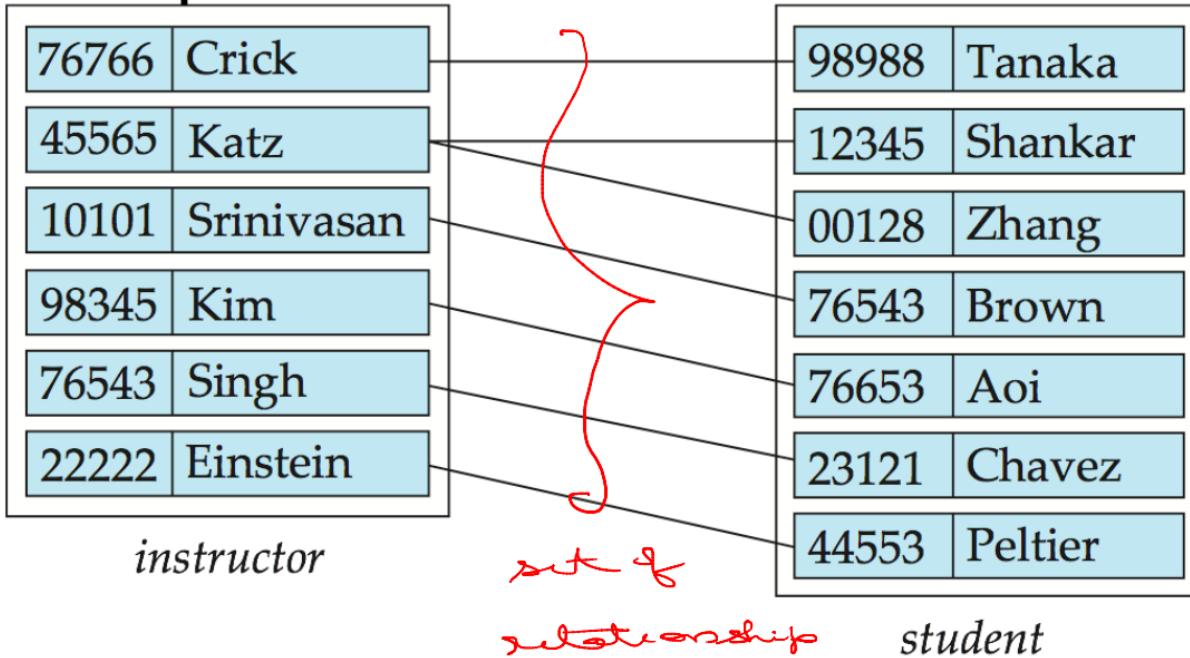
$$\{(e_1, e_2, \dots, e_n) \mid e_1 \in E_1, e_2 \in E_2, \dots, e_n \in E_n\}$$

where  $(e_1, e_2, \dots, e_n)$  is a relationship

- Example:

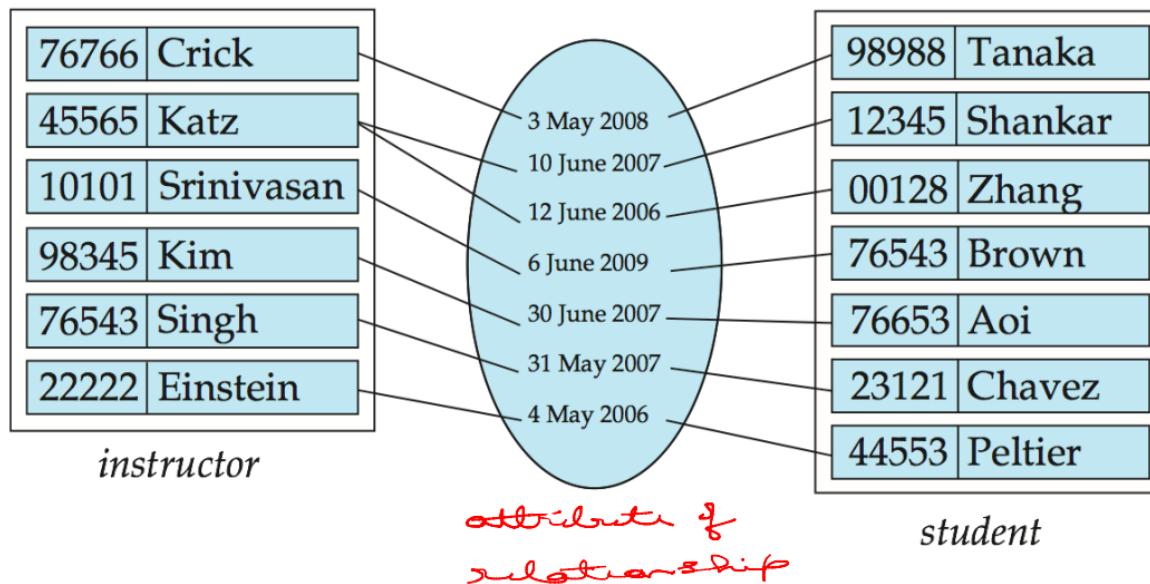
(44553,22222)  $\in$  advisor

# Relationship Set *advisor*



# Relationship Sets (Cont.)

- A **descriptive attribute** can also be **property of a relationship set**.
- For instance, the *advisor* relationship set between entity sets *instructor* and *student* may have the **attribute date** which tracks when the student started being associated with the advisor



# Degree of a Relationship Set

- **binary relationship**



- involve **two entity sets** (or degree two).
- most relationship sets in a database system are binary.

- Relationships between more than two entity sets are rare. **Most relationships are binary**. (More on this later.)



❑ Example: *students* work on research *projects* under the guidance of an *instructor*.

❑ relationship *proj\_guide* is a **ternary relationship** between *instructor*, *student*, and *project*

# Attributes

- An entity is represented by a set of attributes , that is descriptive properties possessed by all members of an entity set.

- Example:

*instructor = (ID, name, street, city, salary )*

*course= (course\_id, title, credits)*

- Domain – the set of permitted values for each attribute

- Attribute types:

- Simple and composite attributes.

- Single-valued and multivalued attributes

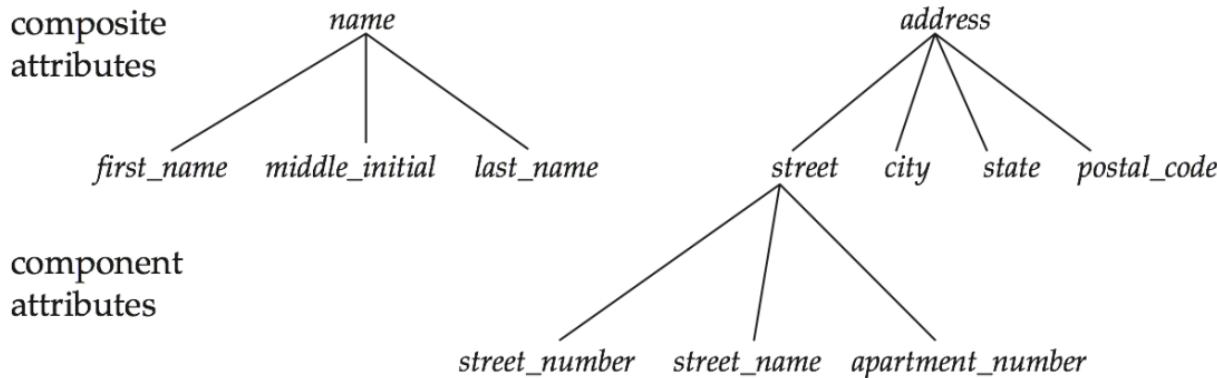
- Example: multivalued attribute: *phone\_numbers*

- Derived attributes

- Can be computed from other attributes

- Example: age, given date\_of\_birth

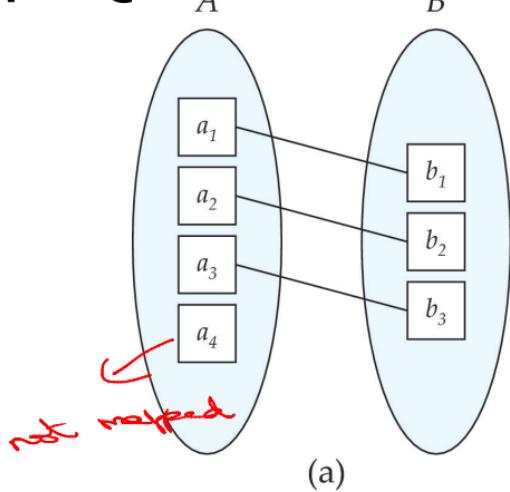
# Composite ~~Attributes~~



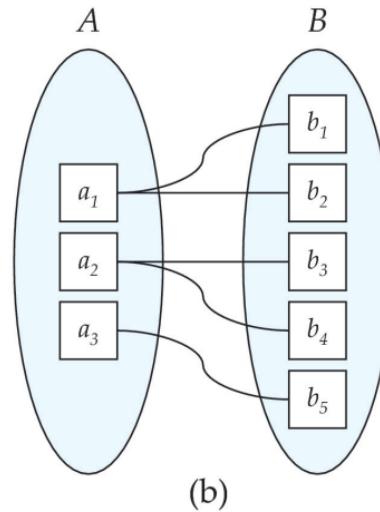
# Mapping Cardinality Constraints

- Express the number of entities to which another entity can be associated via **a relationship set**.
- Most useful in describing binary relationship sets.
- For a binary relationship set the **mapping cardinality** must be one of the following types:
  - One to one
  - One to many
  - Many to one
  - Many to many

# Mapping Cardinalities



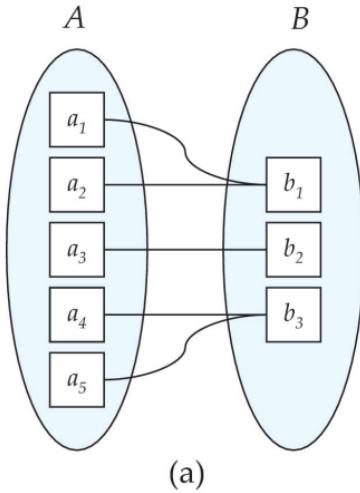
One to one



One to many

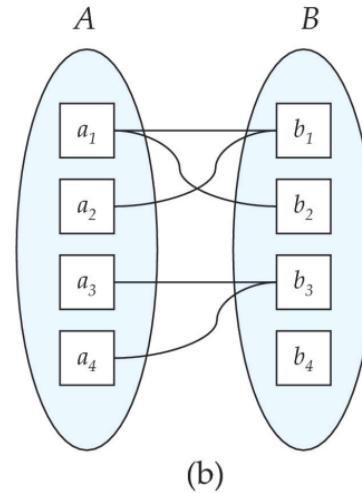
Note: Some elements in  $A$  and  $B$  may not be mapped to any elements in the other set

# Mapping Cardinalities



(a)

Many to  
one



(b)

Many to many

Note: Some elements in A and B may not be mapped to any elements in the other set

# Keys

## Key of an Entity Set

- n A **super key** of an **entity set** is a set of one or more attributes whose values uniquely determine each entity.
- n A **candidate key** of an entity set is a minimal super key
  - / *ID* is candidate key of *instructor*
  - / *course\_id* is candidate key of *course*
- n Although several candidate keys may exist, one of the candidate keys is selected to be the **primary key**.

# Keys for Relationship Sets

- The combination of primary keys of the participating entity sets forms a super key of a relationship set.
  - $(s\_id, i\_id)$  is the super key of *advisor*
  - *NOTE: this means* a pair of entity sets can have at most one relationship in a particular relationship set.
    - Example: if we wish to track multiple meeting dates between a student and her advisor, we cannot assume a relationship for each meeting. We can use a multivalued attribute though
- Must consider the mapping cardinality of the relationship set when deciding what are the candidate keys

# Redundant Attributes

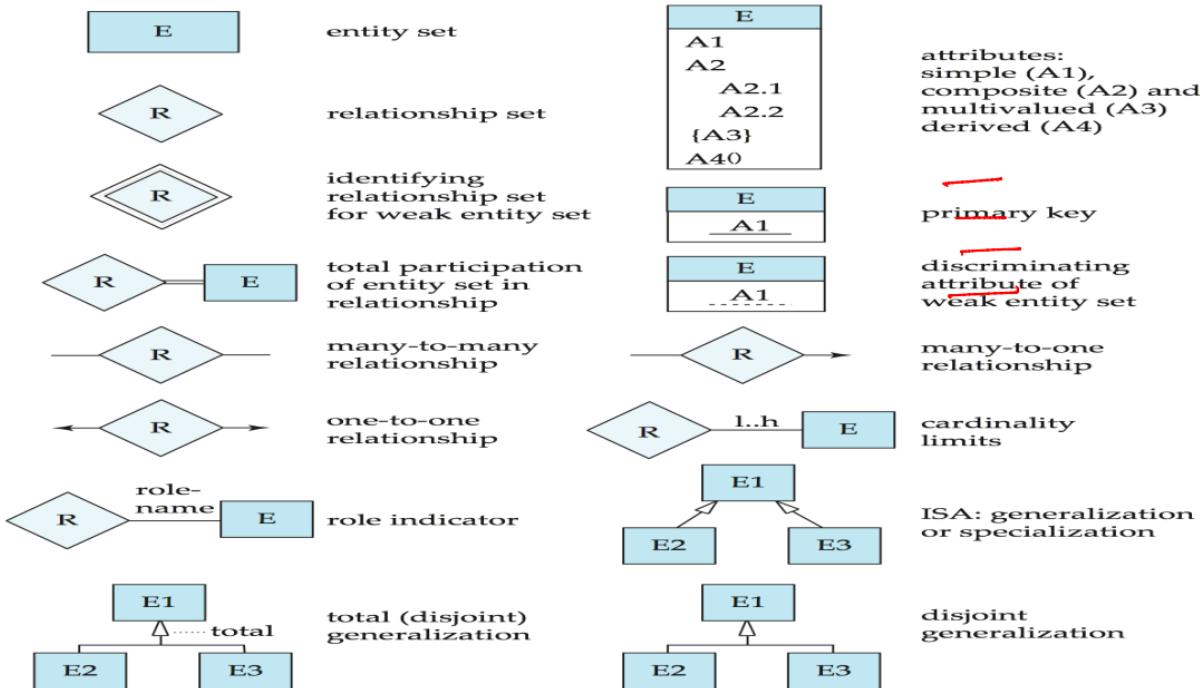
- Suppose we have entity sets
  - *instructor*, with attributes including *dept\_name*
  - *Department*and a relationship
  - *inst\_dept* relating *instructor* and *department*
- Attribute *dept\_name* in entity *instructor* is redundant since there is an explicit relationship *inst\_dept* which relates instructors to departments
  - The attribute replicates information present in the relationship, and should be removed from *instructor*
  - BUT: when converting back to tables, in some cases the attribute gets reintroduced, as we will see.

# ER DIAGRAM

## CHEN'S NOTATION

To be preferably followed

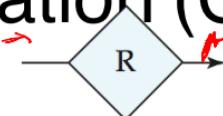
# Symbols Used in E-R Notation



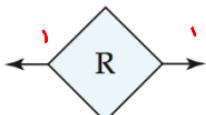
# Symbols Used in E-R Notation (Cont.)



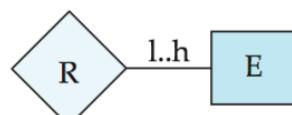
many-to-many  
relationship



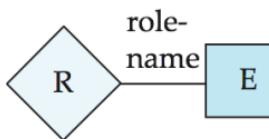
many-to-one  
relationship



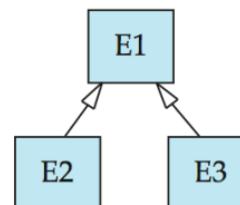
one-to-one  
relationship



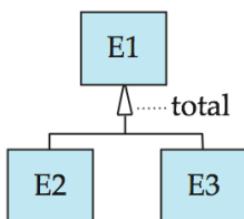
cardinality  
limits



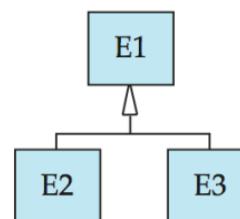
role  
indicator



ISA: generalization  
or specialization



total (disjoint)  
generalization

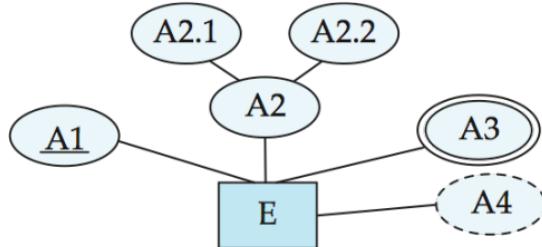


disjoint  
generalization

# Alternative ER Notations

- Chen, IDE1FX, ...

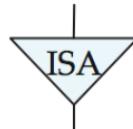
entity set E with  
simple attribute A1,  
composite attribute A2,  
multivalued attribute A3,  
derived attribute A4,  
and primary key A1



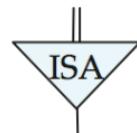
weak entity set



generalization



total  
generalization



# Alternative ER Notations

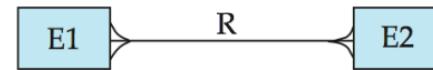
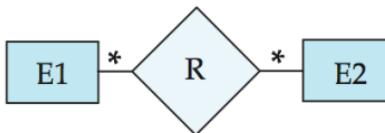
notation)

Chen

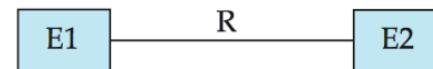
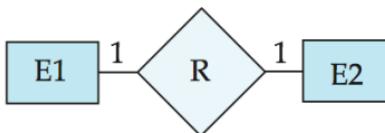
IDE1FX (

Crows feet

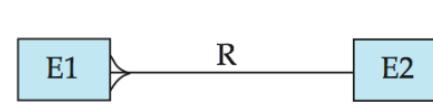
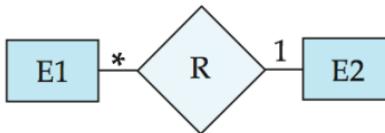
many-to-many  
relationship



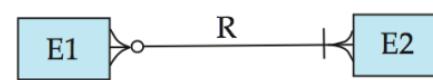
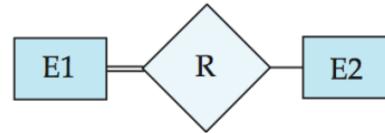
one-to-one  
relationship



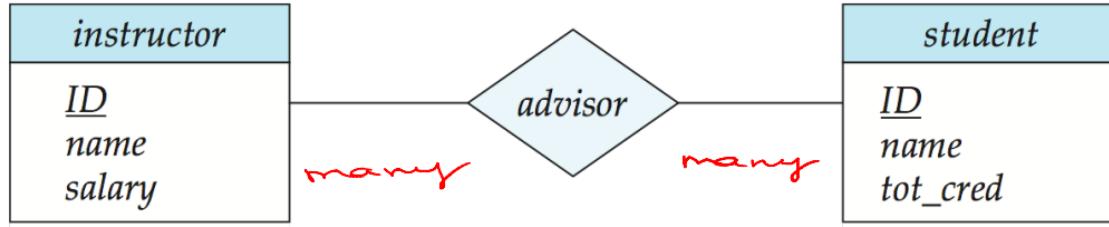
many-to-one  
relationship



participation  
in R: total (E1)  
and partial (E2)



# E-R Diagrams



- n **Rectangles** represent entity sets.
- n **Diamonds** represent relationship sets.
- n Attributes listed inside entity rectangle
- n **Underline** indicates **primary key** attributes

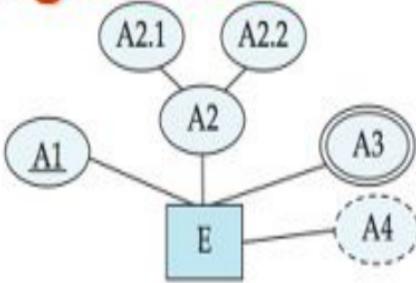
## Entity With Composite, Multivalued, and Derived Attributes

<i>instructor</i>
<u>ID</u>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age</i> ( )

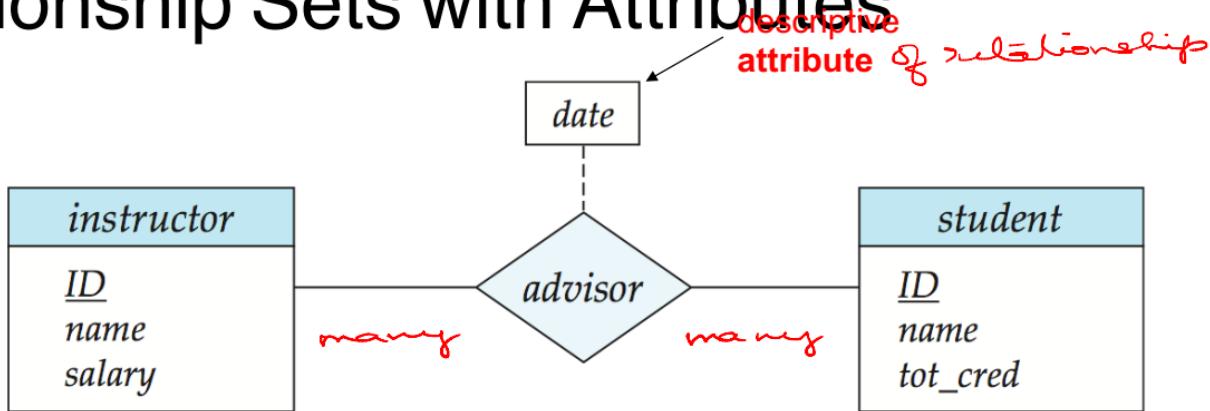
*Note :-*

*Preferable representation*

entity set E with  
simple attribute A1,  
composite attribute A2,  
multivalued attribute A3,  
derived attribute A4,  
and primary key A1



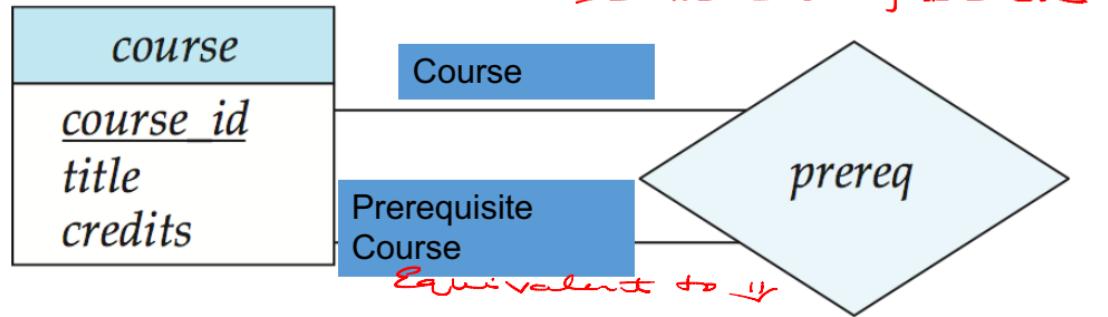
# Relationship Sets with Attributes



# Roles

- Entity sets of a relationship need not be distinct
  - Each occurrence of an entity set plays a “role” in the relationship
- The labels “`course_id`” and “`prereq_id`” are called **roles**.

*Roles important when mapping to  
same Entity is done*



# constraints

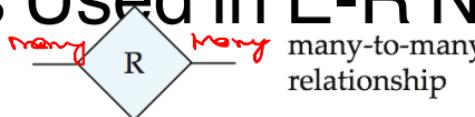
- Cardinality
- participation

# Cardinality Constraints

- We express cardinality constraints by drawing either a directed line (↑), signifying “one,” or an undirected line (—), signifying “many,” between the relationship set and the entity set.

*Note: many means one*

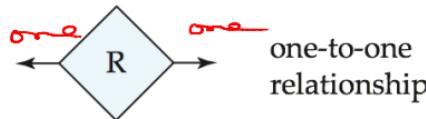
# Symbols Used in E-R Notation (Cont.)



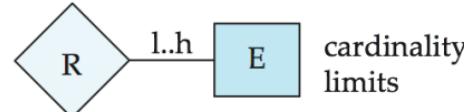
many-to-many  
relationship



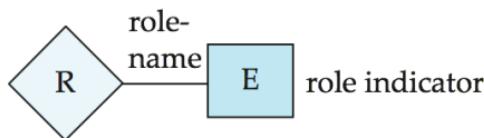
many-to-one  
relationship



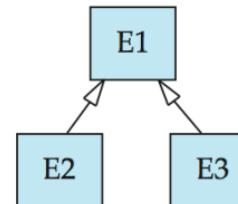
one-to-one  
relationship



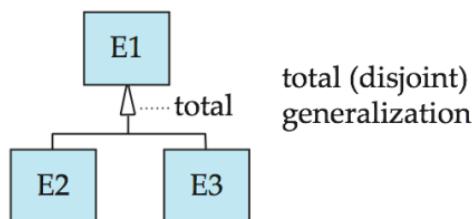
cardinality  
limits



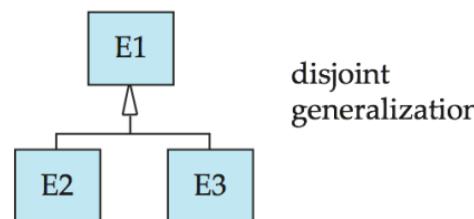
role indicator



ISA: generalization  
or specialization



total (disjoint)  
generalization

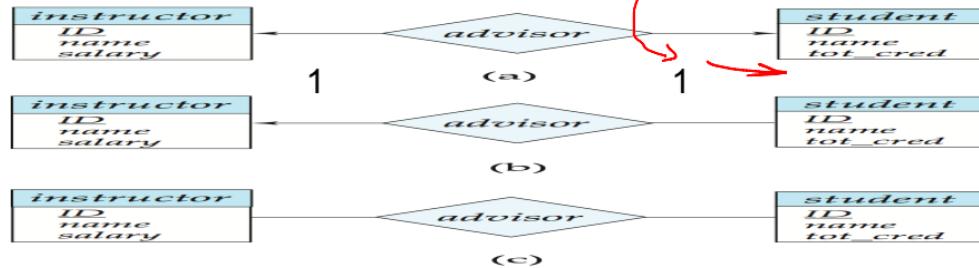


disjoint  
generalization

# One-to-One Relationship

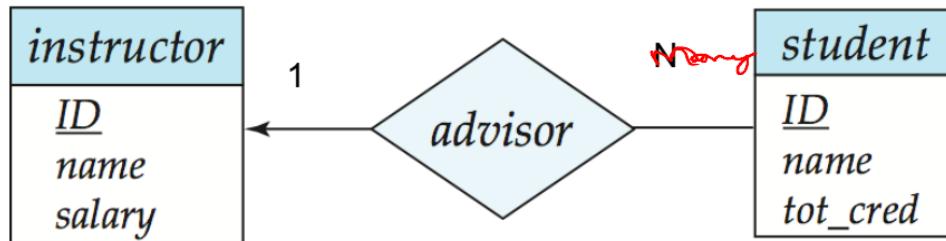
- one-to-one relationship between an *instructor* and a *student*
  - an instructor is associated with **at most one** student via *advisor*
  - and a student is associated with **at most one** instructor via *advisor*

*each instructor associates with  
one student and vice versa*



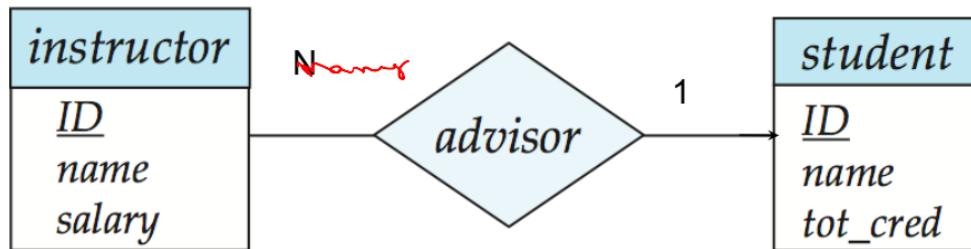
# One-to-Many Relationship

- one-to-many relationship between an *instructor* and a *student*
  - an **instructor** is associated with **several (including 0)** students via *advisor*
  - a student is associated with at most one instructor via *advisor*,



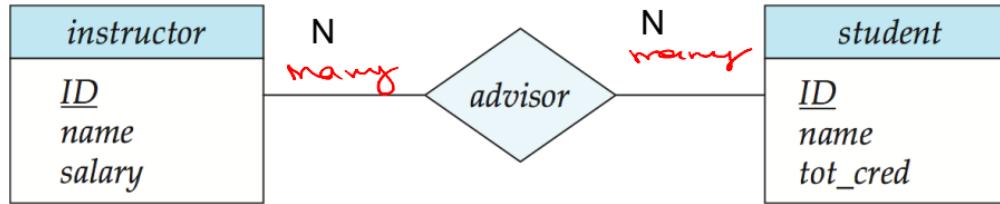
# Many-to-One Relationships

- In a many-to-one relationship between an *instructor* and a *student*,
  - an instructor is associated with **at most one student** via *advisor*,
  - and a student is associated **with several (including 0) instructors** via *advisor*



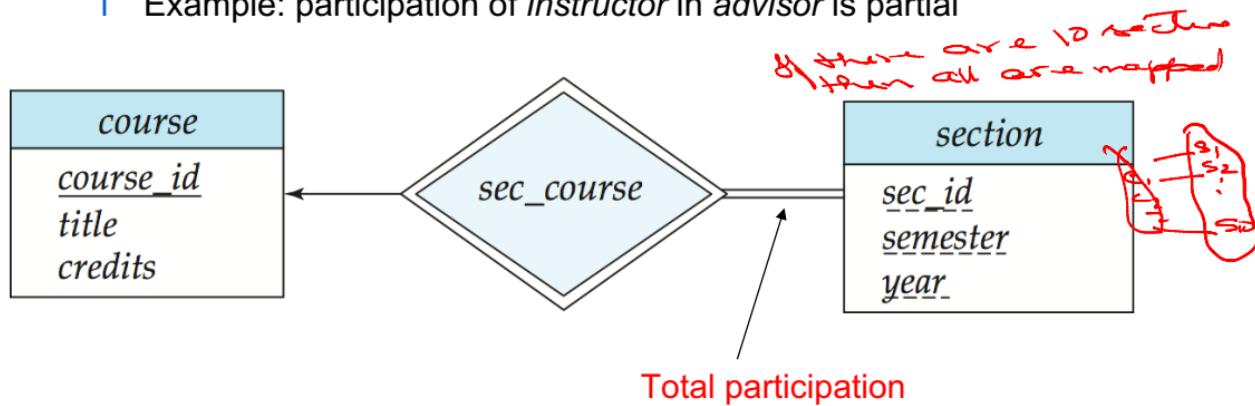
# Many-to-Many Relationship

- An instructor is associated with several (possibly 0) students via *advisor*
- A student is associated with several (possibly 0) instructors via *advisor*



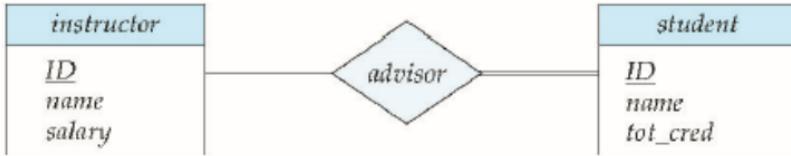
## Participation of an Entity Set in a Relationship Set

- n **Total participation** (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set
  - | E.g., participation of *section* in *sec\_course* is total
    - ?? every section must have an associated course
- n **Partial participation**: some entities may not participate in any relationship in the relationship set
  - | Example: participation of *instructor* in *advisor* is partial



## Total and Partial Participation

- Total participation (indicated by double line): every entity in the entity set participates in at least one relationship in the relationship set

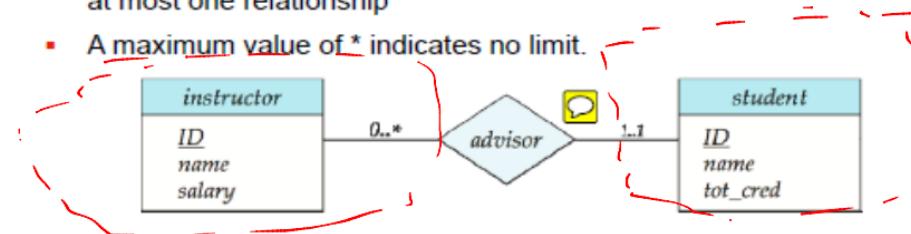


- participation of *student* in *advisor* relation is total
  - every student must have an associated instructor
- Partial participation: some entities may not participate in any relationship in the relationship set
  - Example: participation of *instructor* in *advisor* is partial

*Keywords:-*  
*Every - must -*

## Notation for Expressing More Complex Constraints

- A line may have an associated minimum and maximum cardinality, shown in the form  $l..h$ , where  $l$  is the minimum and  $h$  the maximum cardinality
  - A minimum value of 1 indicates total participation.
  - A maximum value of 1 indicates that the entity participates in at most one relationship
  - A maximum value of \* indicates no limit.

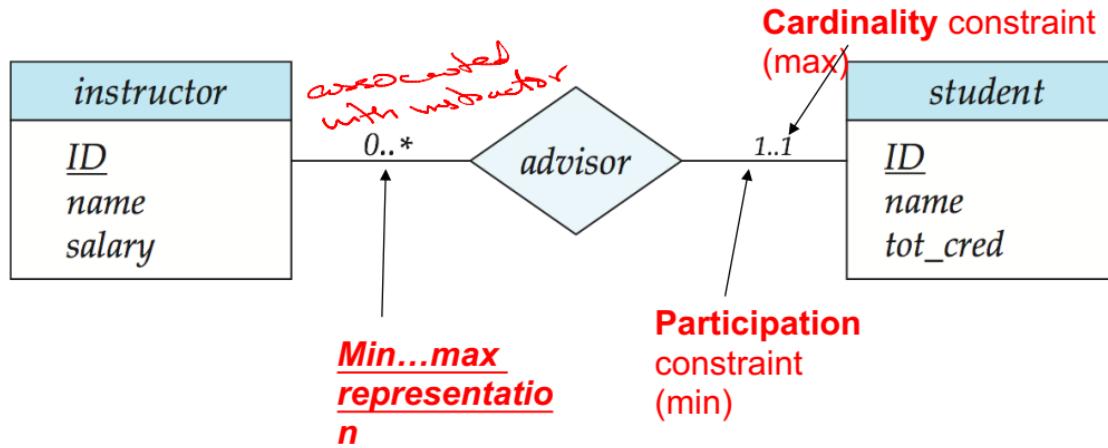


Instructor can advise 0 or more students. A student must have 1 advisor; cannot have multiple advisors

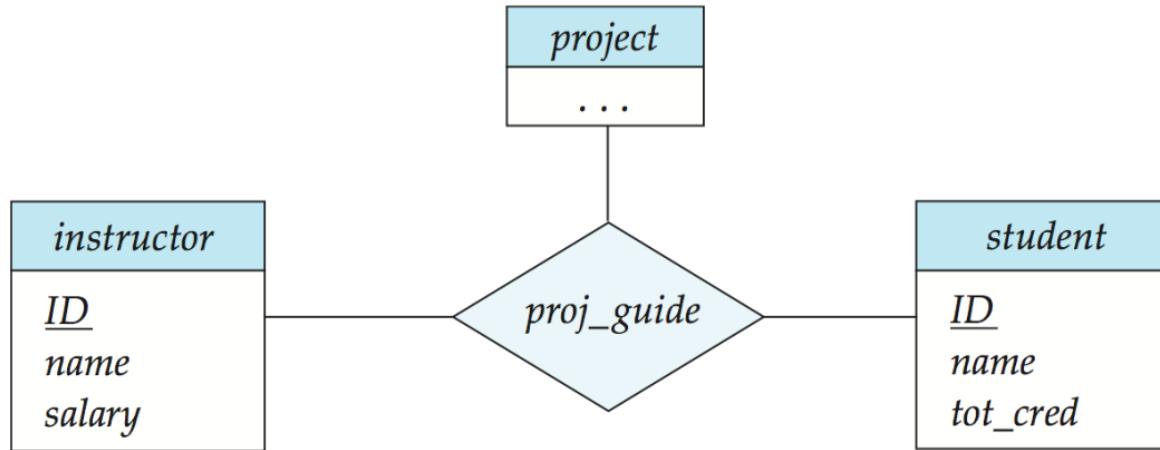
*Note: cardinality associated with the entity close to where it is written*

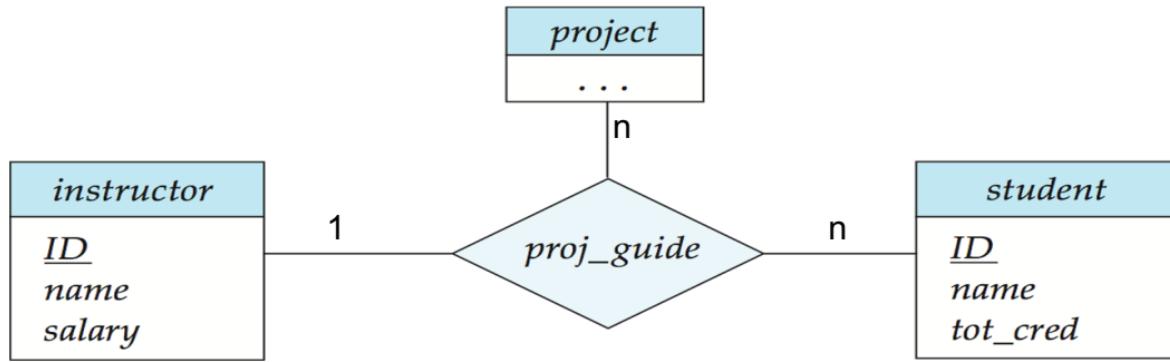
# Alternative Notation for Cardinality Limits

- Cardinality limits can also express participation constraints



# E-R Diagram with a Ternary Relationship

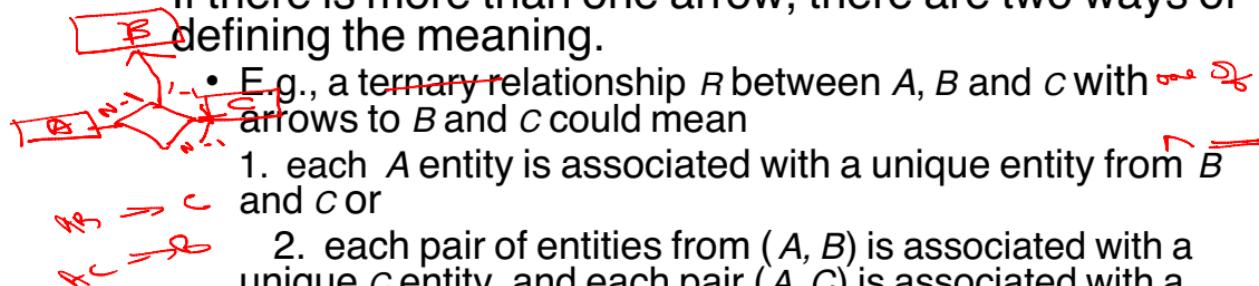




# Cardinality Constraints on Ternary Relationship

Outlaw most one arrow out of a ternary (or greater degree) relationship to indicate a cardinality constraint

- E.g., an arrow from *proj\_guide* to *instructor* indicates each student has at most one guide for a project
- If there is more than one arrow, there are two ways of defining the meaning.



1. each  $A$  entity is associated with a unique entity from  $B$  and  $C$  or
  2. each pair of entities from  $(A, B)$  is associated with a unique  $C$  entity, and each pair  $(A, C)$  is associated with a unique  $B$
- Each alternative has been used in different formalisms
  - To avoid confusion we outlaw more than one arrow

# Weak Entity Sets

- An entity set that does not have a primary key is referred to as a **weak entity set**.
- The existence of a weak entity set depends on the existence of a **identifying entity set**
  - It must relate to the identifying entity set via a total, one-to-many relationship set from the identifying to the weak entity set
  - **Identifying relationship** depicted using a double diamond
- The **discriminator (or partial key) of a weak entity set** is the set of attributes that distinguishes among all the entities of a weak entity set.
- The **primary key of a weak entity** set is formed by the primary key of the strong entity set on which the weak entity set is existence dependent, plus the weak entity set's discriminator.

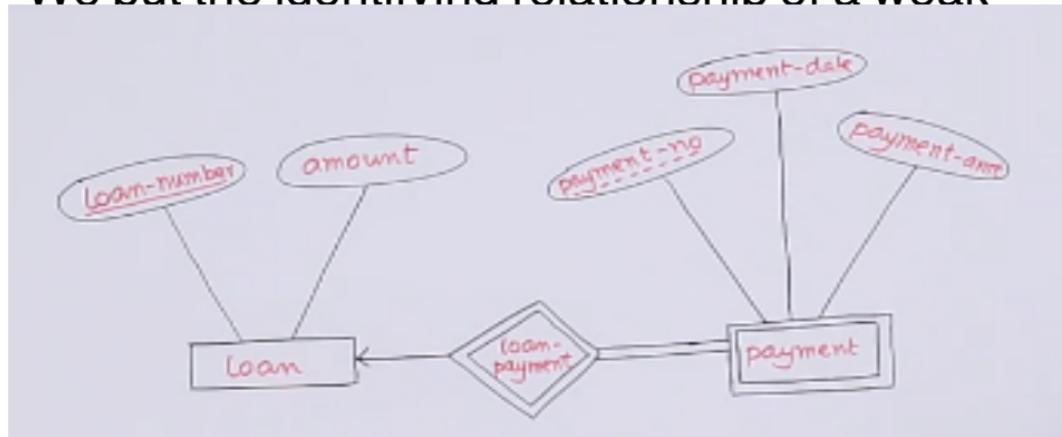
## Expressing Weak Entity Sets

- In E-R diagrams, a weak entity set is depicted via a double rectangle
- We underline the discriminator of a weak entity set with a dashed line
- The relationship set connecting the weak entity set to the identifying strong entity set is depicted by a double diamond
- Primary key for section – (course\_id, sec\_id, semester, year)  
*sections (course\_id, sec\_id, semester, year)*



# Weak Entity Sets (Cont.)

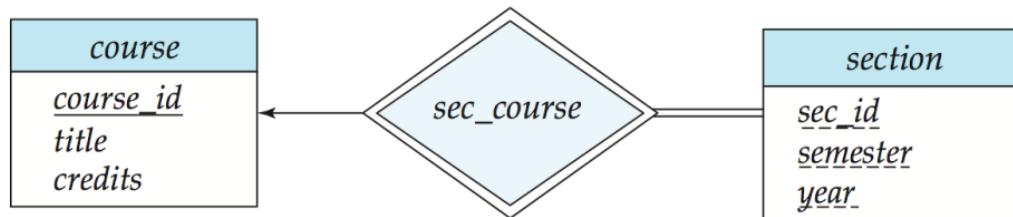
- We underline the **discriminator (or partial key)** of a weak entity set with **a dashed line**.
- We put the identifying relationship of a weak



The **primary key of a weak entity set** is formed by the **primary key of the strong entity set** on which the weak entity set is existence dependent, **plus** the weak entity set's discriminator.

# Weak Entity Sets (Cont.)

- Note: **the primary key of the strong entity set is not explicitly stored with the weak entity set**, since it is implicit in the identifying relationship.
- If *course\_id* were explicitly stored, *section* could be made a strong entity, but then the relationship between *section* and *course* would be duplicated by an implicit relationship defined by the attribute *course\_id* common to *course* and *section*
- Primary key for *section* – (*course\_id*, *sec\_id*, *semester*, *year*)



# Removing redundant attributes in entity sets

- Once the entities and their corresponding attributes are chosen, the relationship sets among the various entities are formed. These relationship sets may result in a situation where attributes in the various entity sets are redundant and need to be removed from the original entity sets.

- the attribute *dept name* in fact gets added to the relation *instructor*, but only if each instructor has at most one associated department.
- If an instructor has more than one associated department, the relationship between instructors and departments is recorded in a separate relation *inst dept*.
- Ex2: sec- timeslot
- And so on

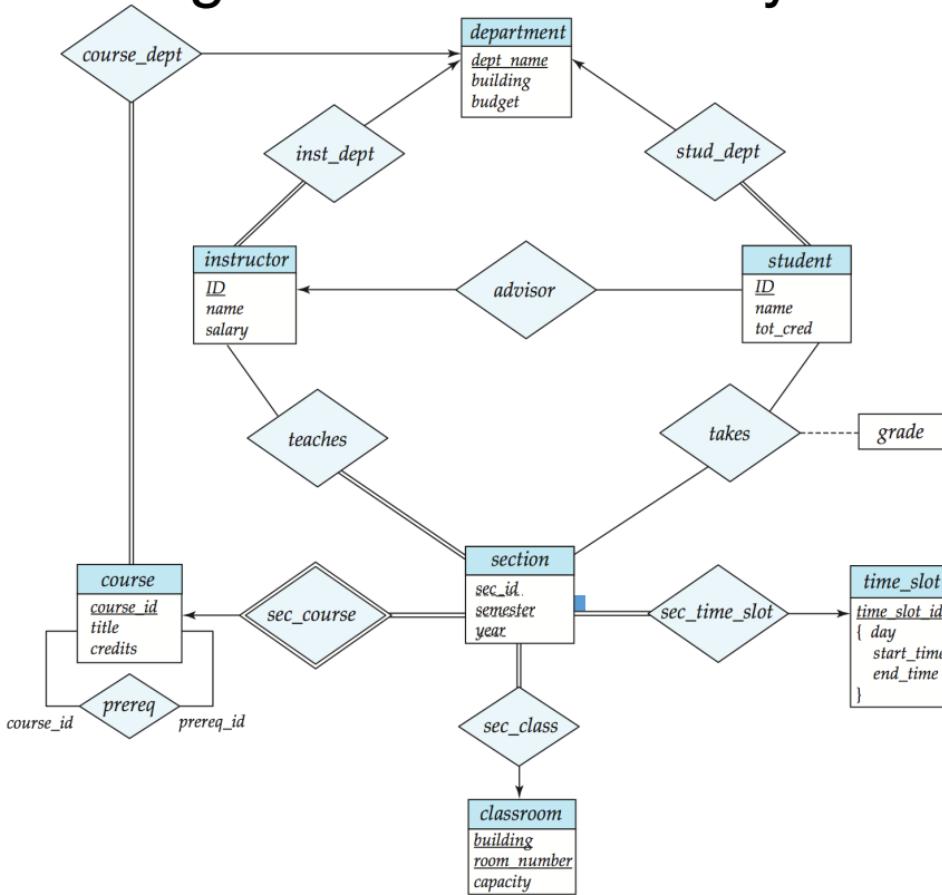
For our university example, we list the entity sets and their attributes below, with primary keys underlined:

- **classroom**: with attributes (*building, room number, capacity*).
- **department**: with attributes (*dept name, building, budget*).
- **course**: with attributes (*course id, title, credits*).
- **instructor**: with attributes (*ID, name, salary*).
- **section**: with attributes (*course id, sec id, semester, year*).
- **student**: with attributes (*ID, name, tot cred*).
- **time slot**: with attributes (*time slot id, {(day, start time, end time) }*).

The relationship sets in our design are listed below:

- **inst dept**: relating instructors with departments.
- **stud dept**: relating students with departments.
- **teaches**: relating instructors with sections.
- **takes**: relating students with sections, with a descriptive attribute *grade*.
- **course dept**: relating courses with departments.
- **sec course**: relating sections with courses.
- **sec class**: relating sections with classrooms.
- **sec time slot**: relating sections with time slots.
- **advisor**: relating students with instructors.
- **prereq**: relating courses with prerequisite courses.

# E-R Diagram for a University Enterprise



# ER diagram

- **E-R diagram** can express the overall logical structure of a database graphically.
- Basic structure
  - cardinality mapping
  - participation
- Weak entity
- Composite
- multivalued

# Reduction to Relational Schemas

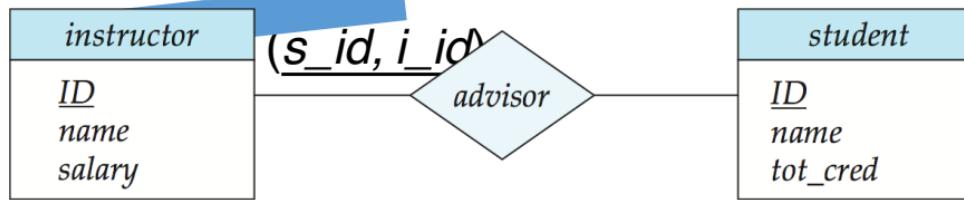
# Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as *relation schemas* that represent the contents of the database.
- A database which conforms to an E-R diagram can be represented by a collection of schemas.
- For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set.
- Each schema has a number of columns (generally corresponding to attributes), which have unique names.

# Representing Relationship Sets

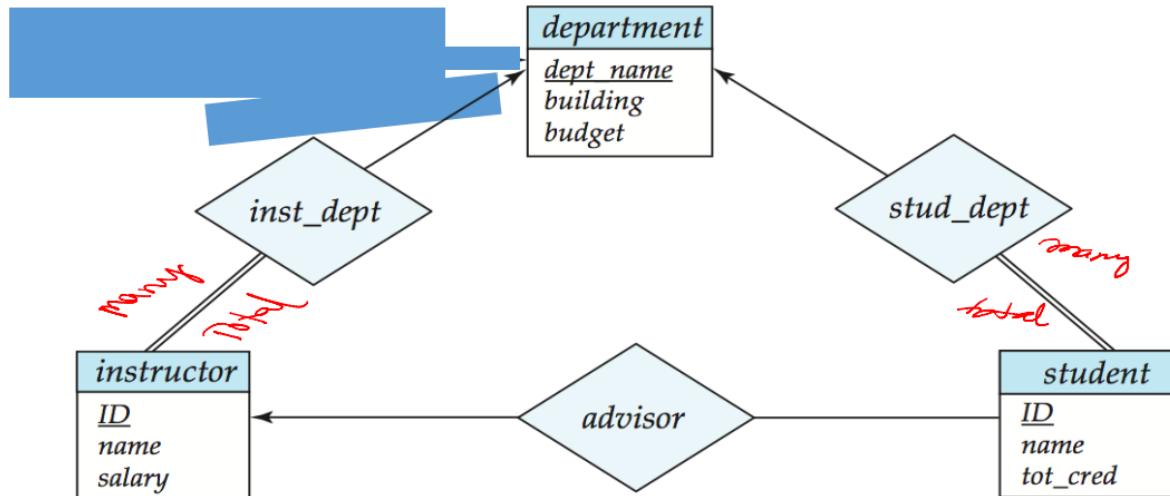
Rule 1:

- A many-to-many relationship set is represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set.
- Example: schema for relationship set *advisor*



## Rule 2

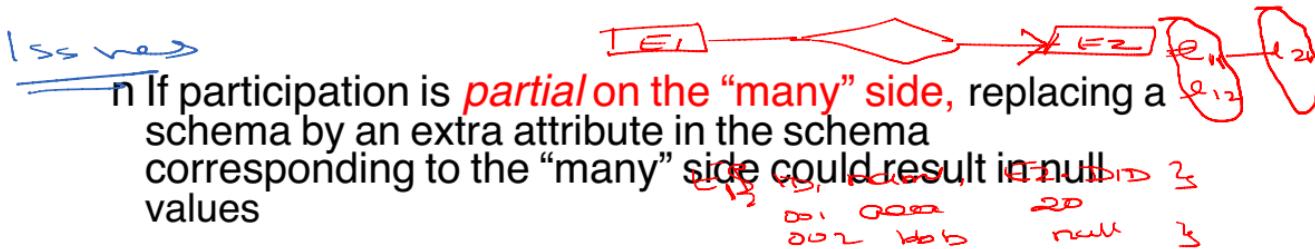
- n Many-to-one and one-to-many relationship sets that are total on the many-side can be represented by adding an extra attribute to the "many" side, containing the primary key of the "one" side
- n If Partial participation: Separate table
- n Example: Instead of creating a schema for relationship set *inst\_dept*, add an attribute *dept\_name* to the schema arising from entity set *instructor*



## Redundancy of Schemas (Cont.)

n For **one-to-one** relationship sets, either side can be chosen to act as the “many” side

- I That is, **extra attribute can be added to either of the tables** corresponding to the two entity sets



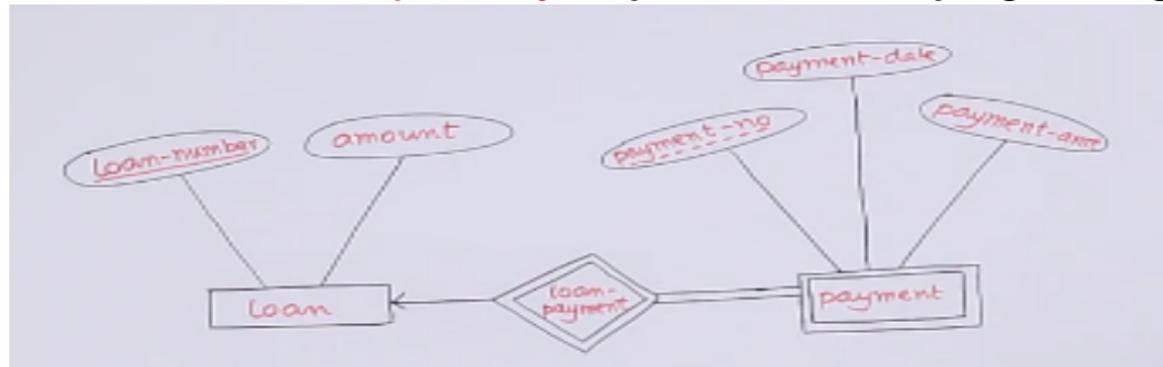
n The **schema corresponding to a relationship set linking a weak entity set to its identifying strong entity set is redundant.**

- I Example: The *section* schema already contains the attributes that would appear in the *sec\_course* schema

## Representing Entity Sets With Simple Attributes

Rule 4:

- A strong entity set reduces to a schema with the same attributes  
*student(ID, name, tot\_cred)*
- A weak entity set becomes a table that includes a column for the primary key of the identifying strong



# Composite and Multivalued Attributes

Rules

<i>instructor</i>
<i>ID</i>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age()</i>

- Composite attributes are flattened out by creating a separate attribute for each component attribute
  - Example: given entity set *instructor* with composite attribute *name* with component attributes *first\_name* and *last\_name* the schema corresponding to the entity set has two attributes *name\_first\_name* and *name\_last\_name*
    - Prefix omitted if there is no ambiguity
- Ignoring multivalued attributes, extended instructor schema is
  - *instructor(ID, first\_name, middle\_initial, last\_name, street\_number, street\_name, apt\_number, city, state, zip\_code, date\_of\_birth)*

# Composite and Multivalued Attributes

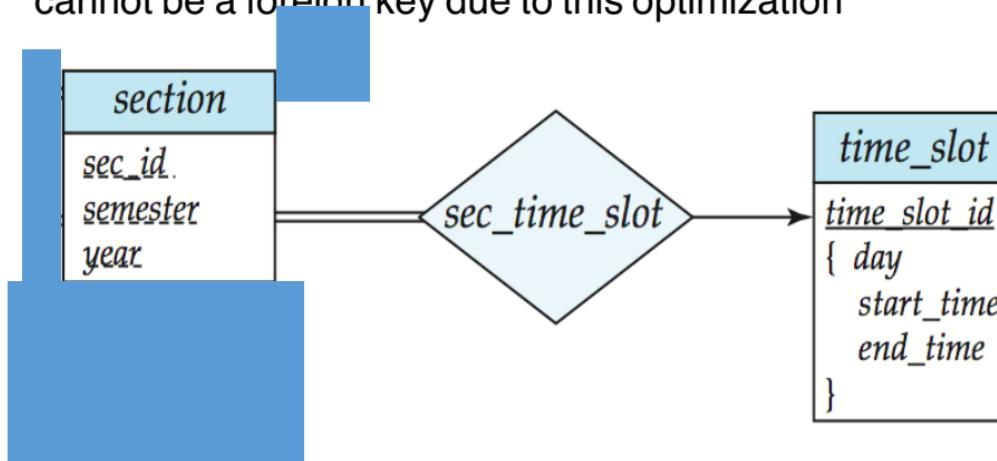
Rule 6

- A **multivalued attribute  $M$**  of an entity  $E$  is \_\_\_\_\_ represented by a separate schema  $EM$ 
  - Schema  $EM$  has attributes corresponding to the primary key of  $E$  and an attribute corresponding to multivalued attribute  $M$
  - Example: Multivalued attribute *phone\_number* of *instructor* is represented by a schema:  
 $inst\_phone = (\underline{ID}, \underline{phone\_number})$
  - Each value of the multivalued attribute maps to a separate tuple of the relation on schema  $EM$ 
    - For example, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples: (22222, 456-7890) and (22222, 123-4567)

# Multivalued Attributes (Cont.)

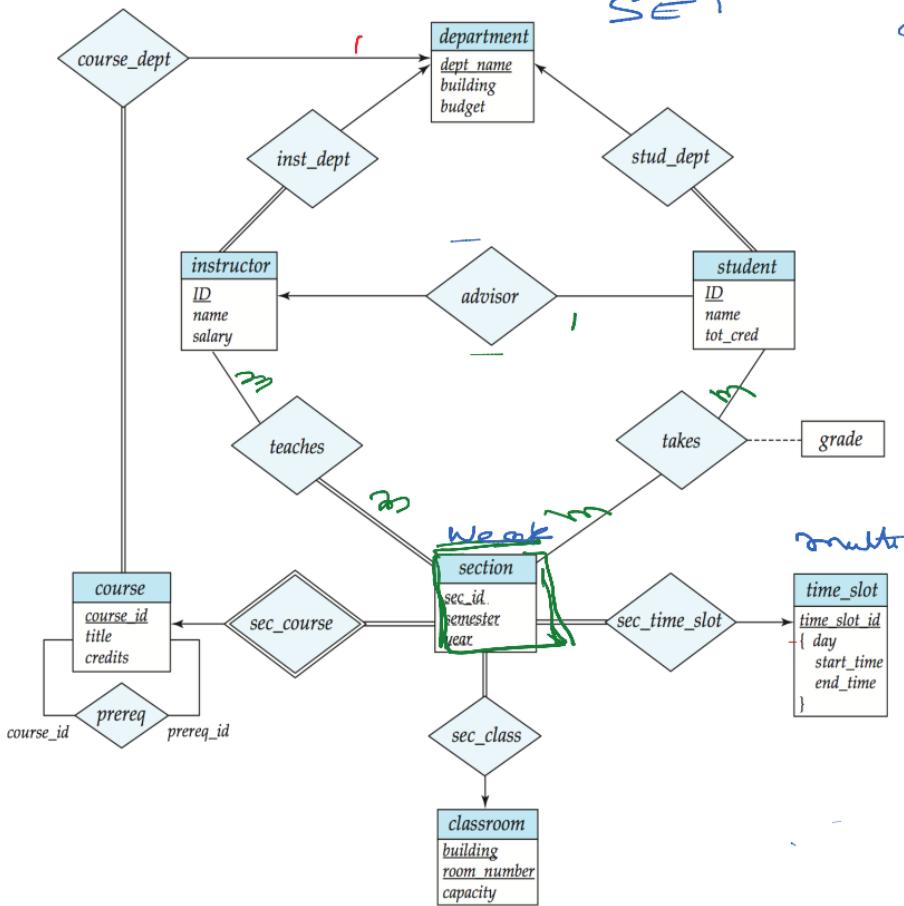
Special case: entity *time\_slot* has only one attribute other than the primary-key attribute, and that attribute is multivalued

- Optimization: Don't create the relation corresponding to the entity, just create the one corresponding to the multivalued attribute
- time\_slot*(*time\_slot\_id*, *day*, *start\_time*, *end\_time*)
- Caveat: *time\_slot* attribute of *section* (from *sec\_time\_slot*) cannot be a foreign key due to this optimization



# Summarizing the rule

- Regular entity: schema with attributes
- Weak entities: schema including strong entity PK
- Composite attribute: schema is extended
- Multivalued attribute: new schema with its individual attribute and PK of the parent table
- Based on degree
  - One to one: separate schema
  - Many to one or one to many: schema for many side include PK of one side
  - Many to many: relationship is converted to schema
- Based on participation:
  - Total participation: on many side add PK of one side
  - Partial participation: separate table



SE 1

SE: dept  
information  
class rooms  
Course  
student

NE: sections

ovv: transient

many-many  
teachers  
takes

dept(dname, building, budget)

student(ID, name, tot-crnd, dname, IID)

instructor(IID, name, salary, dname)

,

,

,

classroom(building, roomno, capacity)

course(courseid, title, credit, dname)

course(courseid, title, credit, dname)

teaches (IID, seid, sem, year, courseid)

takes (grade, ID, seid, sem, year, courseid)

time slot(day, start time, end time)

sections (courseid, seid, sem, year, building,

roomno, IID) ;

Advisor(SID, IID)

Prereq(courseid,

prereqid);

# Construct ER diagram (student work)

1. Each person has at most one e-mail address. Each e-mail address belongs to exactly one person.
2. An apartment is located in a house in a street in a city in a country.
3. Two teams play football against each other. A referee makes sure the rules are followed.
4. Construct an E-R diagram for a hospital with a set of patients and a set of medical doctors. Associate with each patient a log of the various tests and examinations conducted.
5. Consider a company database having each employee works for only 1 department which is managed by one of them. Department has projects running under it and assigned to a team of employees. Further employees are managed by a supervisor. Company is also interested to know the details of employee's dependents.