

PART 2

DATABASE DESIGN

Extended ER Features

features of ER:-

- 1) Specialization & generalization
- 2) Aggregation

Extended E-R Features: Specialization

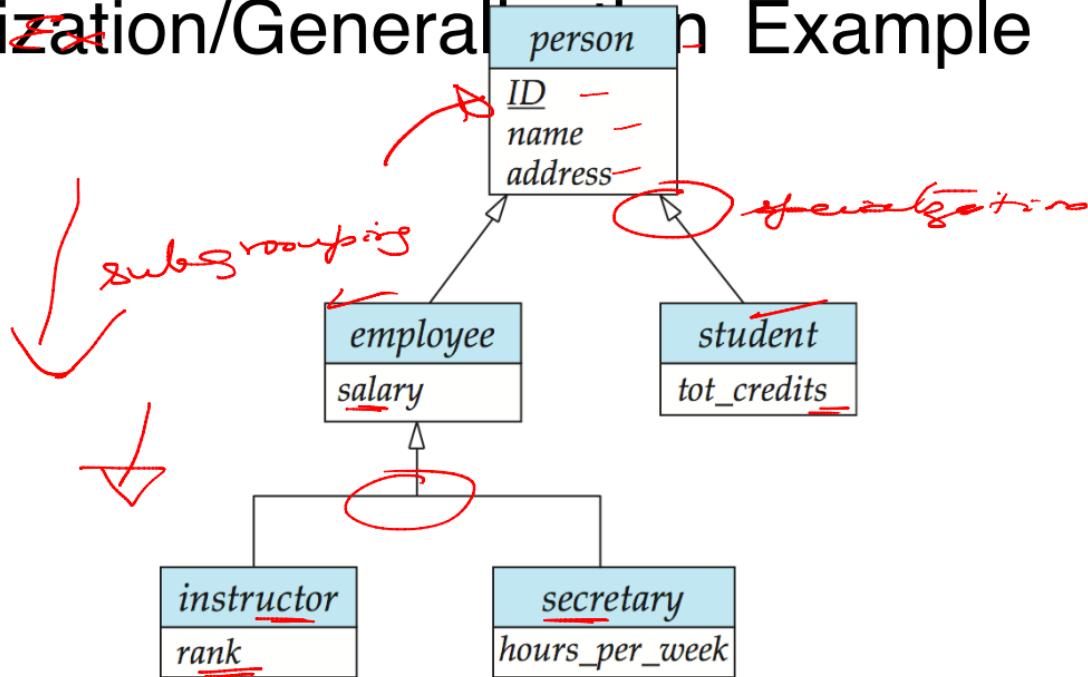
- Top-down design process; we designate subgroupings within an entity set that are distinctive from other entities in the set.
- The process of designating subgroupings within an entity set is called **specialization**.
- These subgroupings become lower-level entity sets that have attributes or participate in relationships that do not apply to the higher-level entity set.
- Depicted by a triangle component labeled ISA (E.g., *instructor* “is a” *person*).



- **Attribute inheritance** – a lower-level entity set inherits all the attributes and relationship participation of the higher-level entity set to which it is linked.

- In terms of an E-R diagram, specialization is depicted by a hollow arrow-head pointing from the specialized entity to the other entity (see Figure 7.21).
- We refer to this relationship as the ISA relationship, which stands for “is a” and represents, for example, that an instructor “is a” employee.
- For an overlapping specialization (as is the case for *student* and *employee* as specializations of *person*), two separate arrows are used.
- For a disjoint specialization (as is the case for *instructor* and *secretary* as specializations of *employee*), a single arrow is used.

Specialization/Generalization Example



Advantage ?

Representing Specialization via Schemas

Method 1:

- Form a schema for the higher-level entity
- Form a schema for each lower-level entity set, include primary key of higher-level entity set and local attributes

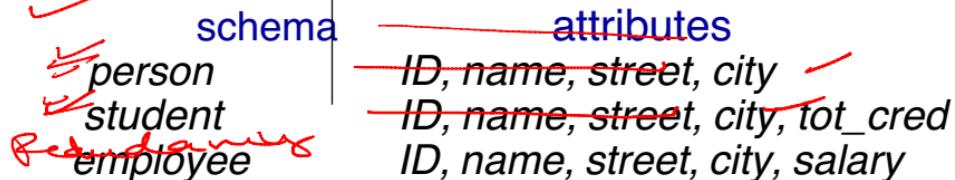


<u>schema</u>	<u>attributes</u>
✓ person	<u>ID</u> , <u>name</u> , <u>street</u> , <u>city</u>
✓ student	<u>ID</u> , <u>tot_cred</u>
✓ employee	<u>ID</u> , <u>salary</u>

- Drawback: getting information about, an *employee* requires accessing **two relations**, the one corresponding to the low-level schema and the one corresponding to the high-level schema

Representing Specialization as Schemas (Cont.)

- Method 2:
 - Form a schema for each entity set with all local and inherited attributes

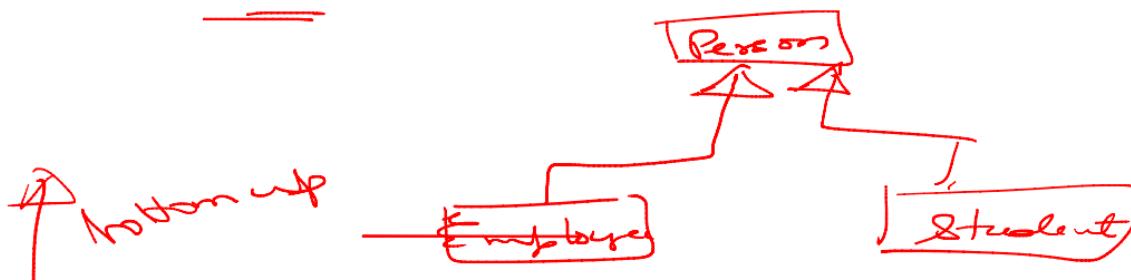


- If specialization is total, the schema for the generalized entity set (*person*) not required to store information
 - Can be defined as a “view” relation containing union of specialization relations
- Drawback: *name, street* and *city* may be stored redundantly for people **who are both students and employees**

Extended ER Features: Generalization

- A bottom-up design process combine a number of entity sets that share the same features into a higher-level entity set.

- Specialization and generalization are simple inversions of each other; they are represented in an E-R diagram in the same way .
- The terms specialization and generalization are used interchangeably.



Specialization and Generalization (Cont.)

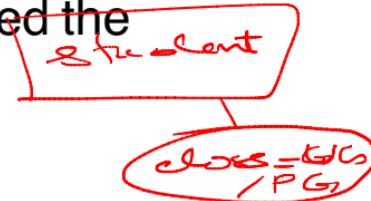
- Can have multiple specializations of an entity set based on different features.
- E.g., ~~permanent_employee vs. temporary_employee~~, in addition to *instructor* vs. *secretary*
- Each particular employee would be
 - a member of one of ~~permanent_employee~~ or *temporary_employee*,
 - and also a member of one of *instructor*, *secretary*
- The ISA relationship also referred to as **superclass - subclass** relationship

Design Constraints on a Specialization/Generalization

Constraint 1: How to group entities?

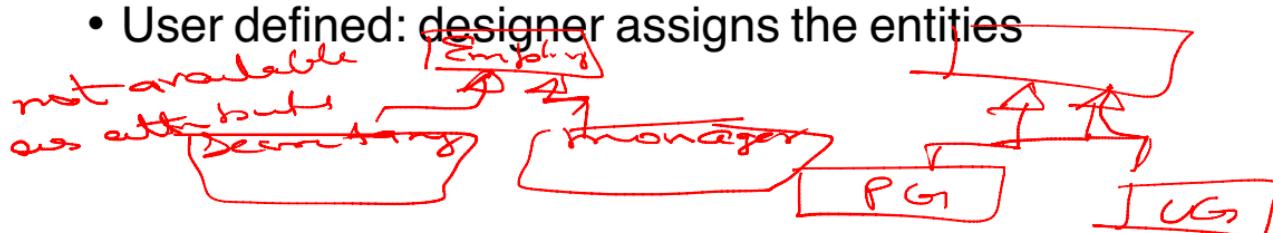
- Conditioned defined: membership evaluated on the basis of if the entity has satisfied the condition

DB



Real world rules

- User defined: designer assigns the entities

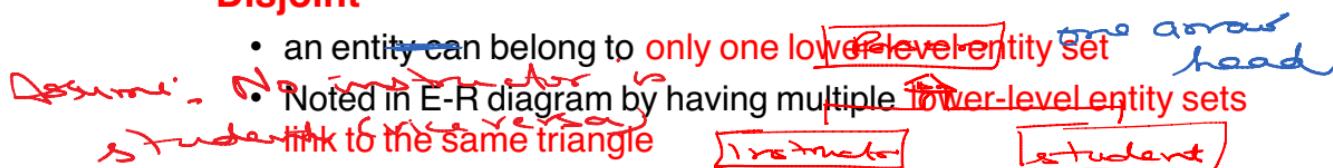


Design Constraints on a Specialization/Generalization

Constraint 2

- Constraint on whether or not entities may belong to more than one lower-level entity set within a single generalization.

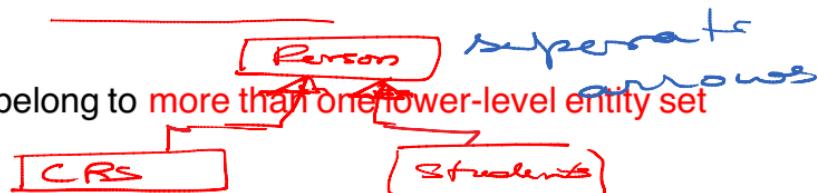
- **Disjoint**



- **Overlapping**

~~Students are
CRS~~

- an entity can belong to more than one lower-level entity set

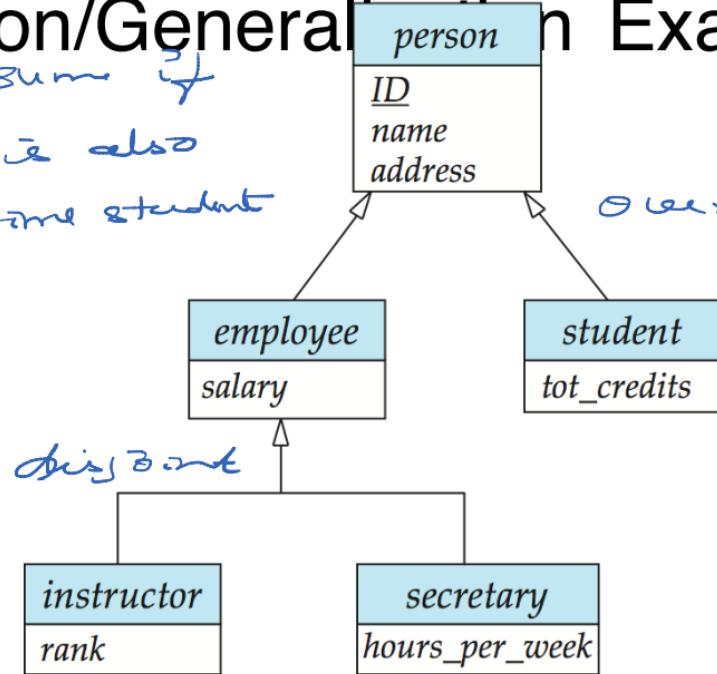


Specialization/Generalization Example

note: Assume

employee is also
a part-time student

overlapping



Advantage ?

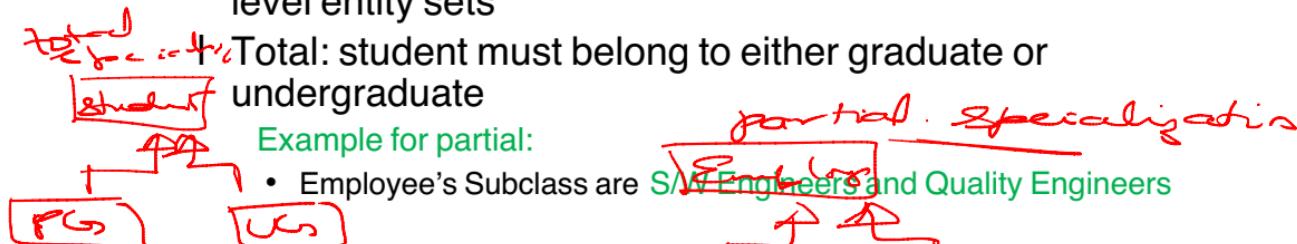
Design Constraints on a Specialization/Generalization (Cont.)

n Completeness constraint -- specifies whether or not an entity in the higher-level entity set **must belong to at least one of the lower-level entity sets** within a generalization.

I **total**: an entity **must belong to one of the** lower-level entity sets

I **partial**: an entity **need not belong to one** of the lower-level entity sets

Total: student must belong to either graduate or undergraduate



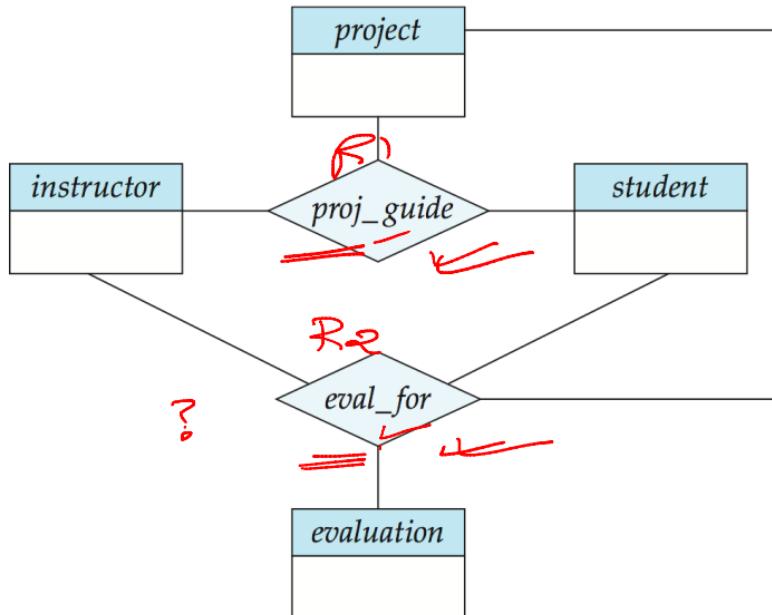
Example for partial:

- Employee's Subclass are SWE Engineers and Quality Engineers

Real world:

Aggregation

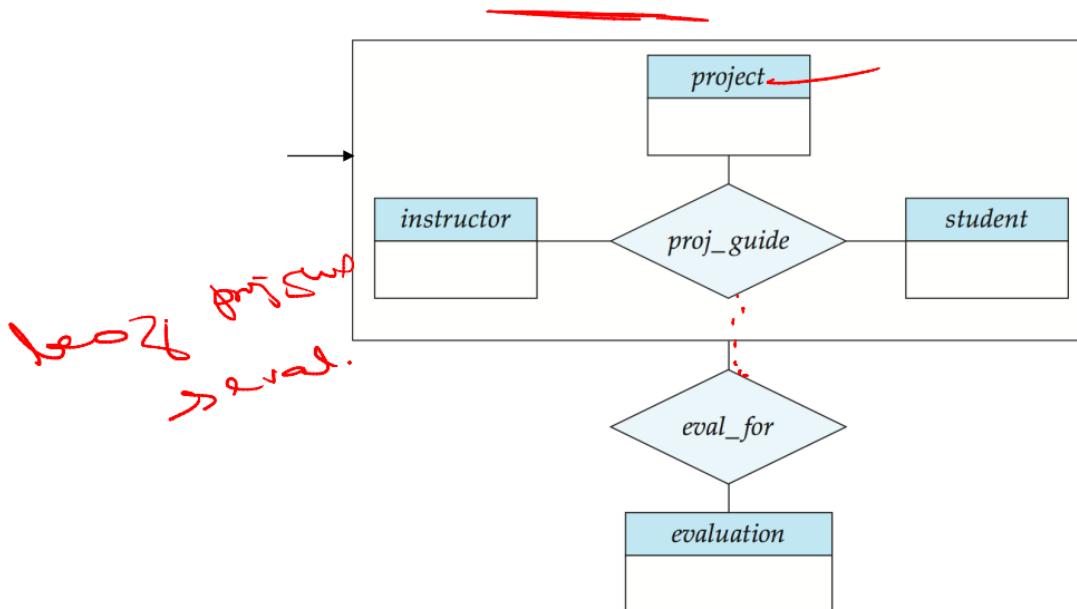
- n Consider the ternary relationship *proj_guide*, which we saw earlier
- n Suppose we want to record evaluations of a student by a guide on a project



Aggregation (Cont.)

- Without introducing redundancy, the following diagram represents:
 - A student is guided by a particular instructor on a particular project
 - A student, instructor, project combination may have an associated evaluation
 - Allows relationships between relationships

Abstract entity



- Aggregation is an abstraction ~~is~~ through which relationships are treated as higher-level entities. Thus the relationship between entities and is treated as if it were an entity.

*represented by bounded rectangle around the entities
and its associated relation*

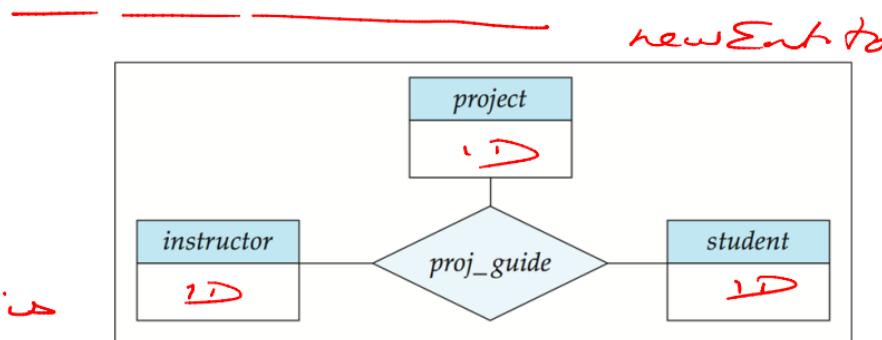
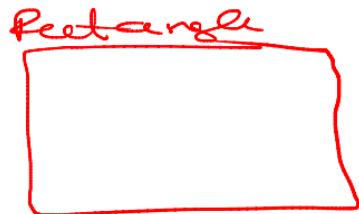
Aggregation (Cont.)

- Relationship sets *eval_for* and *proj_guide* represent overlapping information
 - Every *eval_for* relationship corresponds to a *proj_guide* relationship
 - However, some *proj_guide* relationships may not correspond to any *eval_for* relationships
 - So we can't discard the *proj_guide* relationship
- Eliminate this redundancy via *aggregation*
 - Treat relationship as an *abstract entity*
 - Allows relationships between relationships
 - Abstraction of relationship into new entity

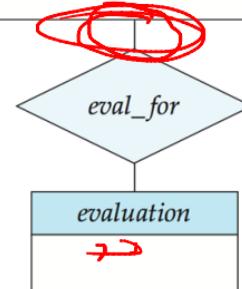
Schemas Corresponding to Aggregation

- n To represent aggregation, create a schema containing
 - | primary key of the aggregated relationship,
 - | the primary key of the associated entity set
 - | any descriptive attributes

eval_for (s_ID, project_id, i_ID, evaluation_id)



*"Relationship is
elevated to an
entity"*

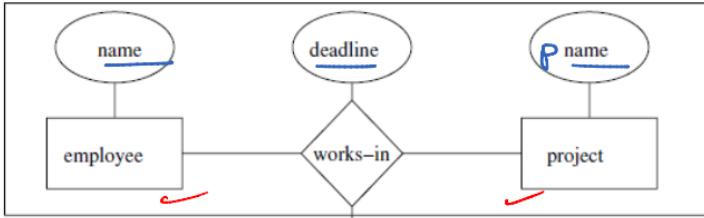


Projects

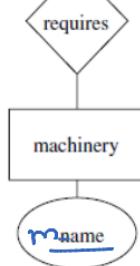
- Employee project us

a particular

binary



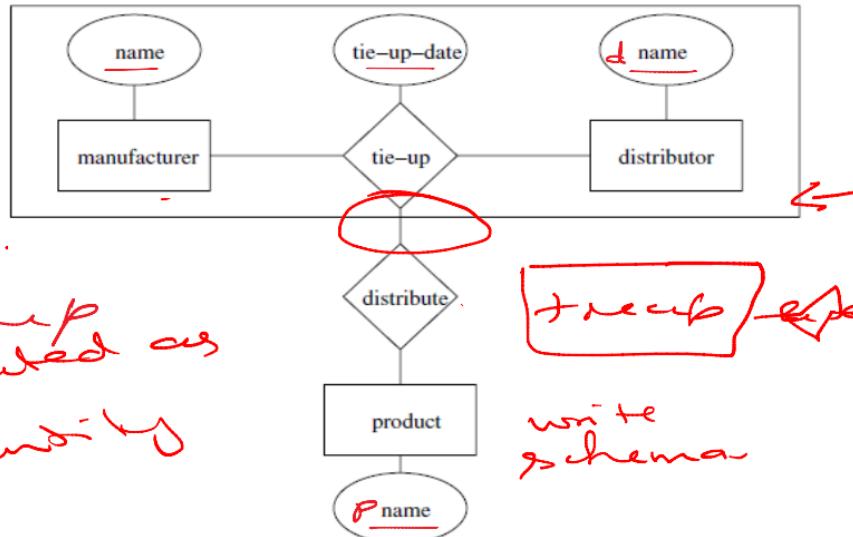
works-in
derived
as entity



entity scheme.

Requires(name, deadline, Pname, m_name) :-

- Manufacturers have tie-ups with distributors to distribute products.
Each tie-up is to be distributed by one distributor.



distribute (name, tieupdate, dname, Pname) ;