



# Chapter 14: Transactions

Database System Concepts, 6th  
Ed.

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Chapter 14: Transactions

- n Transaction Concept
- n Transaction State
- n Concurrent Executions
- n Serializability
- n Recoverability
- n Implementation of Isolation
- n Transaction Definition in SQL
- n Testing for Serializability.



# Transactions

A series of one or more SQL statements that are logically related are termed as a Transaction.

E.g:

1. transaction to transfer \$50 from account A to account B:

1. **read(A)**
2.  $A := A - 50$
3. **write(A)**
4. **read(B)**
5.  $B := B + 50$
6. **write(B)**

*synonym for  
go statement  
to do  
shady purpose*



- n A **transaction** is a **unit** of program execution that accesses and possibly updates various data items. Usually, a transaction is initiated by a user program written in a high-level data-manipulation language
- n A transaction is delimited by statements (or function calls) of the form **begin transaction** and **end transaction**. The transaction consists of all operations executed between the **begin transaction** and **end transaction**.
- n This collection of steps must appear to the user as a single, indivisible unit. This “all-or-none” property is referred to as **atomicity**.



- n the database system must take special actions to ensure that transactions operate properly without interference from concurrently executing database statements. This property is referred to as **isolation**.
- n Even if the system ensures correct execution of a transaction, this serves little purpose if the system subsequently crashes and, as a result, the system “forgets” about the transaction. Thus, a transaction’s actions must persist across crashes. This property is referred to as **durability**.
- n if a transaction is run atomically in isolation starting from a consistent database, the database must again be consistent at the end of the transaction. This property is referred to as **consistency**.

$$\text{Transfer } A \Rightarrow B = \sum(A - B)$$



# ACID Properties

b c t

we require that the database system maintain the following properties of the transactions:

**Atomicity.** Either all operations of the transaction are properly reflected in the database or none are.

*single unit*

**Consistency.** Execution of a transaction in isolation preserves the consistency of the database.

**Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions.

Intermediate transaction results must be hidden from other concurrently executed transactions.

That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.

**Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



# Example of Fund Transfer

Transaction to transfer \$50 from account A to account B:

1. **read(A)**
2.  $A := A - 50$
3. **write(A)**
4. **read(B)**
5.  $B := B + 50$
- write(B)**

*Explain properties with an example*

self study

## Atomicity requirement

if the transaction fails after step 3 and before step 6, money will be “lost” leading to an inconsistent database state

Failure could be due to software(R.Error in Query) or hardware(S/W crash)

the system should ensure that updates of a partially executed transaction are not reflected in the database

**Durability requirement** — once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the transaction must persist even if there are software or hardware failures.



# Example of Fund Transfer (Cont.)

Transaction to transfer \$50 from account A to account B:

1. **read(A)**
2.  $A := A - 50$
3. **write(A)**
4. **read(B)**
5.  $B := B + 50$
6. **write(B)**

**Consistency requirement** in above example:

the sum of A and B is unchanged by the execution of the transaction

In general, consistency requirements include

Explicitly specified integrity constraints such as primary keys and foreign keys

**Implicit integrity constraints**

e.g. sum of balances of all accounts, minus sum of loan amounts must equal  
value of cash-in-hand

A transaction must see a consistent database.

During transaction execution the database may be temporarily inconsistent.

When the transaction completes successfully the database must be consistent

Erroneous transaction logic can lead to inconsistency



# Example of Fund Transfer (Cont.)

- n **Isolation requirement** — if between steps 3 and 6, another transaction T2 is allowed to access the partially updated database, it will see an inconsistent database (the sum  $A + B$  will be less than it should be).

T1

1. **read(A)**
2.  $A := A - 50$
3. **write(A)**
4. **read(B)**
5.  $B := B + 50$
6. **write(B)**

T2

**read(A), read(B), print(A+B)**

- n Isolation can be ensured trivially by running transactions **serially** that is, one after the other.
- n However, executing **multiple transactions concurrently has significant benefits**, as we will see later.





# Storage Structure

- n To understand how to ensure the atomicity and durability properties of a transaction, we must gain a better understanding of how the various data items in the database may be stored and accessed.
- n **Volatile storage:** Information residing in volatile storage does not usually survive system crashes. Examples: **main memory** and **cache memory**.
- n **Nonvolatile storage:** Information residing in nonvolatile storage survives system crashes. Ex: **secondary storage** devices such as magnetic disk and flash storage, used for online storage, and tertiary storage devices such as optical media, and magnetic tapes, used for archival storage.
- n **Stable storage:** Information residing in stable storage is *never lost*.  
*Extremely unlikely to lose data- replicate information*





- n For a transaction to be durable, its changes need to be written to stable storage.
- n Similarly, for a transaction to be atomic, log records need to be written to stable storage before any changes are made to the database on disk.
- n Clearly, the degree to which a system ensures durability and atomicity depends on how stable its implementation of stable storage really is.



# Transaction Atomicity and Durability

Technology

A transaction may not always complete its execution successfully. Such a transaction is termed as **aborted**.

Fails

Once the changes caused by an aborted transaction have been undone, we say that the transaction has been **rolled back**.

It is part of the responsibility of the recovery scheme to manage transaction aborts. This is done typically by maintaining a **log**.

Maintaining a log provides the possibility of redoing a modification to ensure atomicity and durability as well as the possibility of undoing a modification to ensure atomicity in case of a failure during transaction execution.

A transaction that completes its execution successfully is said to be **committed**. The only way to undo the effects of a committed transaction is to execute a **compensating transaction**.



# Transaction State

- n **Active** – the initial state; the transaction stays in this state while it is executing
- n **Partially committed** – after the final statement has been executed.
- n **Failed** -- after the discovery that normal execution can no longer proceed.
- n **Aborted** – after the transaction has been rolled back and the database restored to its state prior to the start of the transaction.  
Two options after it has been aborted:
  - | restart the transaction
  - 4 | can be done only if no internal logical error
  - | kill the transaction
- n **Committed** – after successful completion.

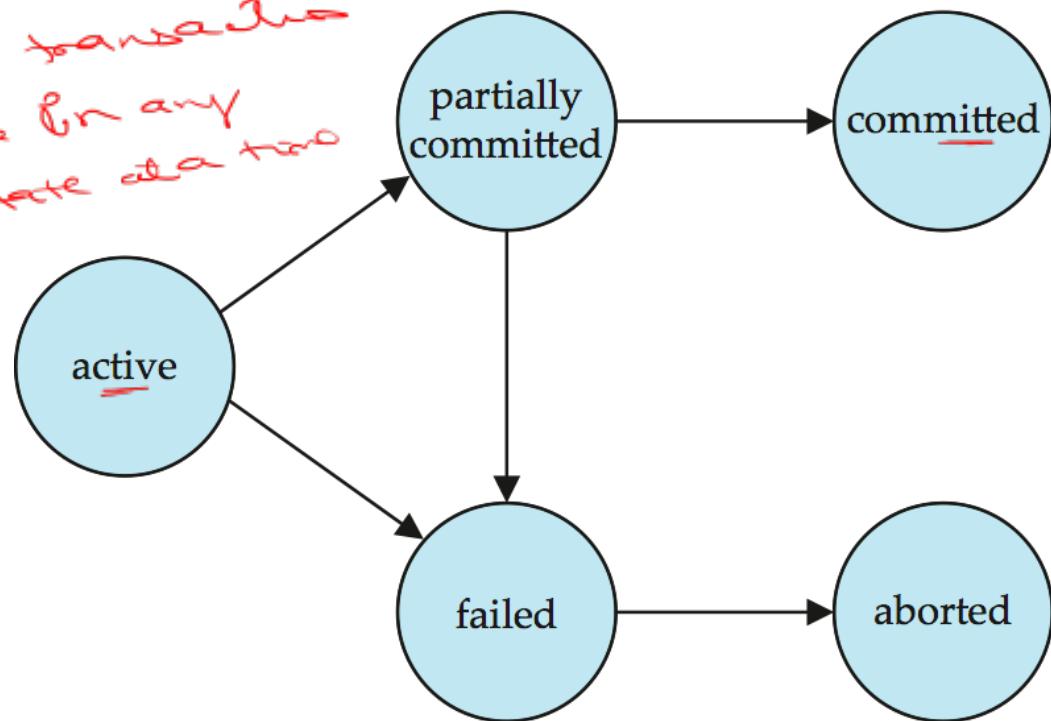




# Transaction model

(\*)

single transaction  
can be in any  
one state at a time



ANSWER



# Transaction State (Cont.)

- n The system has two options if the transaction is failed:
  - n It can **restart the transaction, but only if the transaction was aborted as a result of some hardware or software error that was not created through the internal logic of the transaction.**
  - n A restarted transaction is considered to be a new transaction.
  - n It can **kill the transaction**. It usually does so because of some internal logical error that can be corrected only by rewriting the application program, or because the input was bad, or because the desired data were not found in the database.





# Transaction Concept

## COMMIT and ROLLBACK commands

n A transaction begins with the first executable SQL statement after a commit, rollback or connection made to the database.

— —

n A transaction can be closed by using a commit or a rollback statement.

—

n **COMMIT** ends the current transaction and makes permanent changes made during the transaction

n **ROLLBACK** ends the transaction but undoes any changes made during the transaction

—

n Syntax: COMMIT; ROLLBACK ;

—





## Example for usage of **Commit**, **Rollback** and **SAVEPOINT**

n Withdraw an amount 2,000 and deposit 10,000 for all the accounts. If the sum of balance of all accounts exceeds 2,00,000 then **undo** the deposit just made.

```
Declare  
    Total_bal numeric;  
Begin  
    Update account set balance=balance-2000;  
    Savepoint deposit;  
    Update account set balance=balance+10000;  
    Select sum(balance) into total_bal from account;  
    If total_bal > 200000 then  
        Rollback to savepoint deposit;  
    End if;  
    Commit;  
End;
```

The code is annotated with red hand-drawn markings:

- A red bracket groups the first two statements: `Update account set balance=balance-2000;` and `Savepoint deposit;`.
- A red bracket groups the next three statements: `Update account set balance=balance+10000;`, `Select sum(balance) into total_bal from account;`, and `If total_bal > 200000 then`.
- A red bracket groups the statement `Rollback to savepoint deposit;`.
- A red bracket groups the statements `End if;` and `Commit;`.
- A red bracket groups the entire transaction block: `Update account set balance=balance-2000;`, `Savepoint deposit;`, the three-line block, and `Rollback to savepoint deposit;`.
- A red bracket labeled "undo" is placed to the right of the `End if;` and `Commit;` statements, indicating that the transaction can be rolled back to the savepoint.

n **SAVEPOINT** marks and saves the current point in the processing of a transaction. When a savepoint is used with a **ROLLBACK** statement, parts of transaction can be undone.



# Transaction isolation

- n Ensuring consistency in spite of concurrent execution of transactions requires extra work; it is far easier to insist that transactions run **serially—that is, one at a time, each** starting only after the previous one has completed.
- n Multiple transactions are allowed to run concurrently in the system.
  - | Advantages are:
    - | **increased processor and disk utilization**, leading to better transaction throughput
      - 4 E.g. one transaction can be using the CPU while another is reading from or writing to the disk
    - | **reduced average response time/waiting time** for transactions: short transactions need not wait behind long ones.
- n **Concurrency control schemes** – mechanisms to achieve isolation
  - | that is, to control the interaction among the concurrent transactions in order to prevent them from destroying the **consistency of the database**.



# Why Concurrency Control is needed



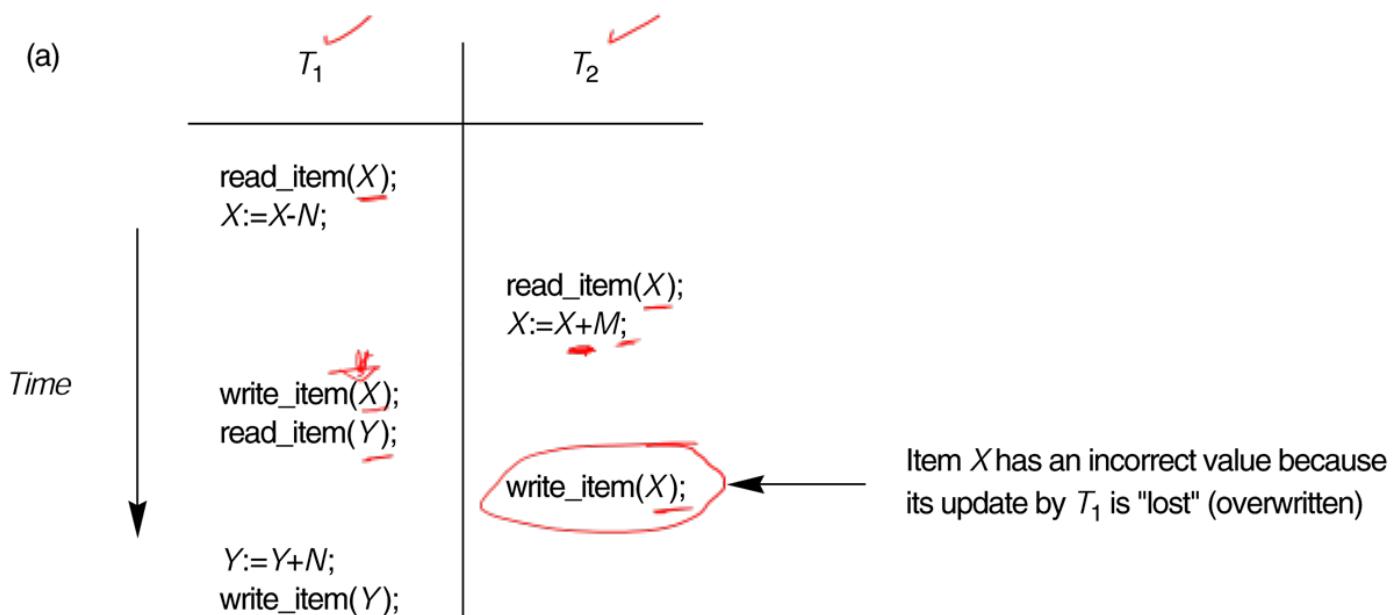
- n Problems occur when concurrent transactions execute in an **uncontrolled** manner:
  - | The Lost Update
  - | The Temporary Update (or Dirty Read)
  - | The Incorrect Summary
  - | Unrepeatable Read:
    - 4 A transaction T1 may read a given value. If another transaction later updates that value and T1 reads that value again, then T1 will see a different value.





# The Lost Update Problem

- This occurs when **two** transactions that access the **same** database items have their operations interleaved in a way that makes the value of some database item incorrect.

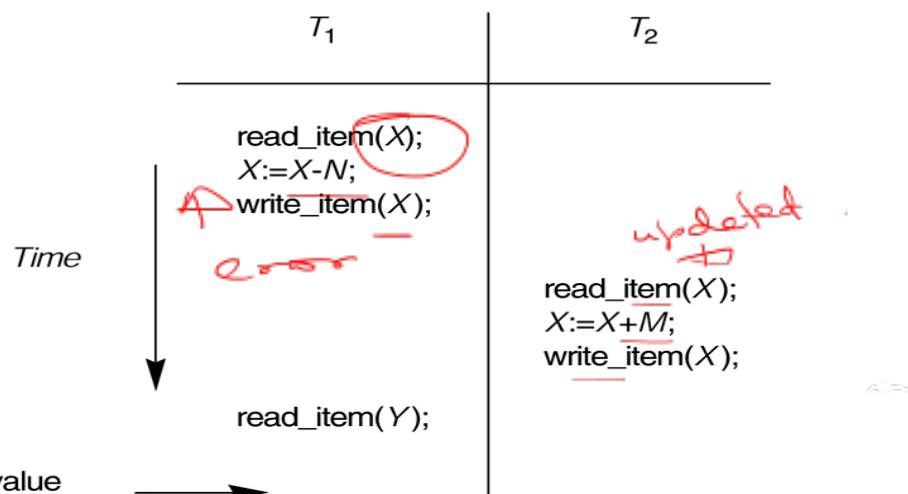




# The Temporary Update Problem (Dirty Read)

- | This occurs when one transaction updates a database item and then the transaction fails for some reason.
- | The updated item is accessed by another transaction before it is changed back to its original value.

(b)



Transaction  $T_1$  fails and must change the value of  $X$  back to its old value; meanwhile  $T_2$  has read the "temporary" incorrect value of  $X$ .



# Incorrect Summary Problem

- n If one transaction is calculating an aggregate summary function on a number of records while other transactions are updating some of these records, the aggregate function may calculate some values before they are updated and others after they are updated.



# Incorrect Summary Problem

**Figure 17.3**

Some problems that occur when concurrent execution is uncontrolled. (a) The lost update problem. (b) The temporary update problem. (c) The incorrect summary problem.

(c)

$T_1$	$T_3$
<pre>sum := 0; read_item(A); sum := sum + A;  •  read_item(X); X := X - N; write_item(X);  read_item(Y); Y := Y + N; write_item(Y);</pre>	<pre>sum := 0; read_item(A); sum := sum + A;  •  read_item(X); sum := sum + X; read_item(Y); sum := sum + Y;</pre> <p><math>\blackleftarrow T_3</math> reads <math>X</math> after <math>N</math> is subtracted and reads <math>Y</math> before <math>N</math> is added; a <u>wrong</u> summary is the result (off by <math>N</math>).</p>



# Schedules

- n **Schedule** – a sequences of instructions that specify the chronological order in which instructions of concurrent transactions are executed
  - | a schedule for a set of transactions must consist of all instructions of those transactions
  - | must preserve the order in which the instructions appear in each individual transaction.
- n A transaction that successfully completes its execution will have a commit instructions as the last statement
  - | by default transaction assumed to execute commit instruction as its last step
- n A transaction that fails to successfully complete its execution will have an abort instruction as the last statement





# Schedules

- n A **schedule**  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is an ordering of **all** the operations in these transactions subject to the constraint that:
  - | for each transaction  $T_i$ , the operations of  $T_i$  in  $S$  must appear in the **same order** as they do in  $T_i$ .
- 4 Note, however, that operations from other transactions  $T_j$  can be interleaved with the operations of  $T_i$  in  $S$ .

- n Example: Given
    - |  $T_1 = R_1(Q) W_1(Q)$  &  $T_2 = R_2(Q) W_2(Q)$
    - | a schedule:  $R_1(Q) R_2(Q) W_1(Q) W_2(Q)$  *order w.r.t. each transaction*
    - | not a schedule:  $W_1(Q) R_1(Q) R_2(Q) W_2(Q)$
- ~~not in specified order~~



9/11/2020

# Schedule 1

- n Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- n A **serial schedule** in which  $T_1$  is followed by  $T_2$  :

*Correct schedule*

$\checkmark T_1$	$T_2 \checkmark$
read ( $A$ )	
$A := A - 50$	
write ( $A$ )	
read ( $B$ )	
$B := B + 50$	
write ( $B$ )	
commit	
	read ( $A$ )
	$temp := A * 0.1$
	$A := A - temp$
	write ( $A$ )
	read ( $B$ )
	$B := B + temp$
	write ( $B$ )
	commit



# Schedule 2

- A **serial schedule** where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	



n

## Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is not a serial schedule (It is CONCURRENT Schedule), but it is equivalent to Schedule 1.

~~serial~~

	$T_1$	$T_2$	
	read ( $A$ ) $A := A - 50$ write ( $A$ )	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ )	
	read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $B$ ) $B := B + temp$ write ( $B$ ) commit	

In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.

ANSWER



# Schedule 4

- n The following concurrent schedule does not preserve the value of  $(A + B)$ .  
It has Lost Update problems.

$T_1$	$T_2$
<p>read (<math>A</math>) <math>A := A - 50</math></p> <p><i>Lost update Doesn't sum up</i></p>	<p>read (<math>A</math>) <math>temp := A * 0.1</math> <math>A := A - temp</math> write (<math>A</math>) read (<math>B</math>)</p>
<p>write (<math>A</math>) read (<math>B</math>) <math>B := B + 50</math> write (<math>B</math>) commit</p>	<p><math>B := B + temp</math> write (<math>B</math>) commit</p>



- n We can ensure consistency of the database under concurrent execution by making sure that any schedule that is executed has the same effect as a schedule that could have occurred without any concurrent execution.
- n That is, the schedule should, in some sense, be equivalent to a serial schedule. Such schedules are called serializable schedules.
- n serial schedules are serializable, but if steps of multiple transactions are interleaved, it is harder to determine whether a schedule is serializable.
- n Since transactions are programs, it is difficult to determine exactly what operations a transaction performs and how operations of various transactions interact.
- n For this reason, we shall not consider the various types of operations that a transaction can perform on a data item, but instead consider only two operations: read and write.



# Serializability

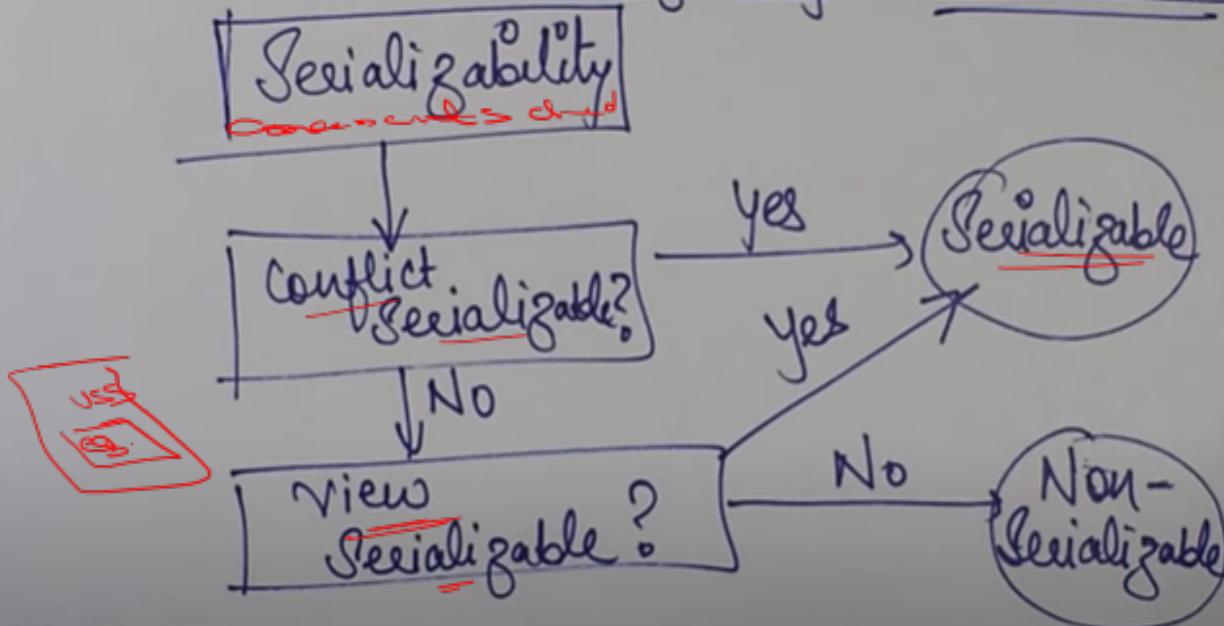
- n **Basic Assumption** – Each transaction preserves database consistency.
- n Thus serial execution of a set of transactions preserves database consistency.
- n A concurrent schedule is **serializable** if it is **equivalent** to a serial schedule.

- 
- 1. **conflict serializability**
  - 2. **view serializability**





## Serializability : A Broader View





# ***Simplified view of transactions***

- | We ignore operations other than **read** and **write** instructions
- | We assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
- | Our simplified schedules consist of only **read** and **write** instructions.



# Conflict Serializability

## Conflicting Instructions

- n Instructions  $l_i$  and  $l_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict if and only if** there exists some item  $Q$  accessed by both  $l_i$  and  $l_j$ , and at least one of these instructions wrote  $Q$ .



1.  $l_i = \text{read}(Q)$ ,  $l_j = \text{read}(Q)$ .  $l_i$  and  $l_j$  don't conflict.
2.  $l_i = \text{read}(Q)$ ,  $l_j = \text{write}(Q)$ . They conflict.
3.  $l_i = \text{write}(Q)$ ,  $l_j = \text{read}(Q)$ . They conflict
4.  $l_i = \text{write}(Q)$ ,  $l_j = \text{write}(Q)$ . They conflict



- n Intuitively, a conflict between  $l_i$  and  $l_j$  forces a (logical) **temporal order between them**.

- | If  $l_i$  and  $l_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.





# Conflict Serializability

*Concurrent schedule*

- n If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are conflict equivalent.
- n We say that a schedule  $S$  is conflict Serializable if it is conflict equivalent to a serial schedule





# Conflict Serializability (Cont.)

- n Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions.
- n Therefore Schedule 3 is conflict serializable.

$T_1$	$T_2$	$T_1$	$T_2$
1 read (A) 2 write (A) 3 read (B) 4 write (B) 5 read (B) 6 write (B)	7 read (A) 8 write (A) 9 read (B) 10 write (B)	1 read (A) 2 write (A) 3 read (B) 4 write (B)	5 read (A) 6 write (A) 7 read (B) 8 write (B)

*non conflicting instructions swapped*

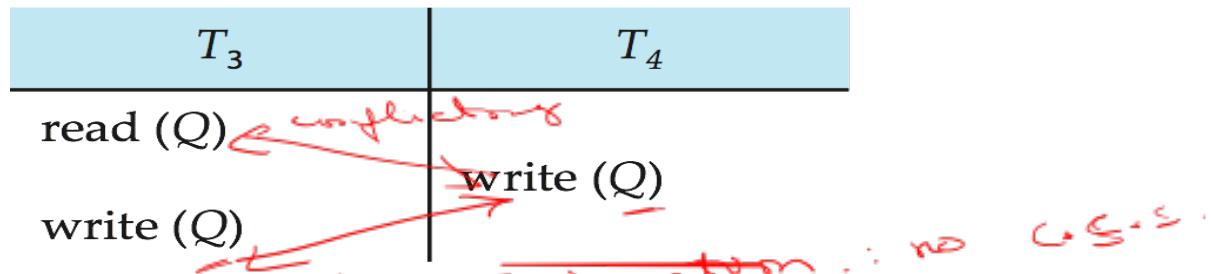
*conflict serial*

Schedule 3                              Schedule 6

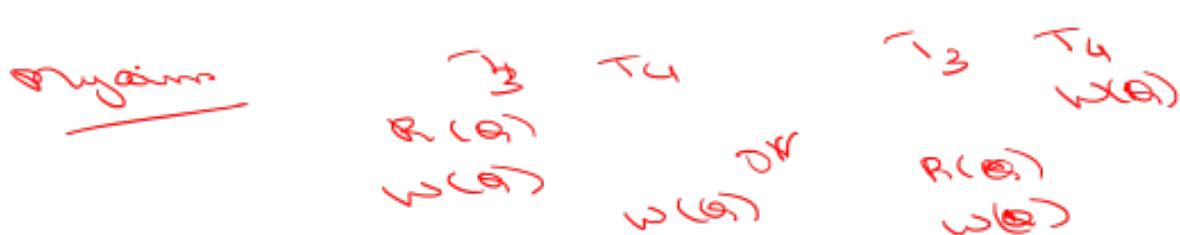


# Conflict Serializability (Cont.)

- n Example of a schedule that is **not conflict serializable**:



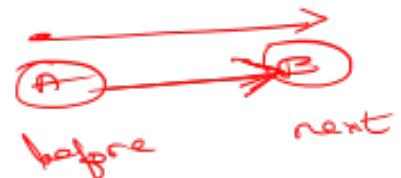
- n We are **unable to swap instructions** in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .





## Check Conflict Serializable, using Precedence Graph

- n We now present a simple and efficient method for determining conflict serializability of a schedule. Consider a schedule  $S$ . We construct a directed graph, called a **precedence graph**, from  $S$ .
- n This graph consists of a pair  $G = (V, E)$ , where  $V$  is a set of vertices and  $E$  is a set of edges. The set of vertices consists of all the transactions participating in the schedule. The set of edges consists of all edges
  - n  $T_i \rightarrow T_j$  for which one of three conditions holds:
    1.  $T_i$  executes write( $Q$ ) before  $T_j$  executes read( $Q$ ).
    2.  $T_i$  executes read( $Q$ ) before  $T_j$  executes write( $Q$ ).
    3.  $T_i$  executes write( $Q$ ) before  $T_j$  executes write( $Q$ ).



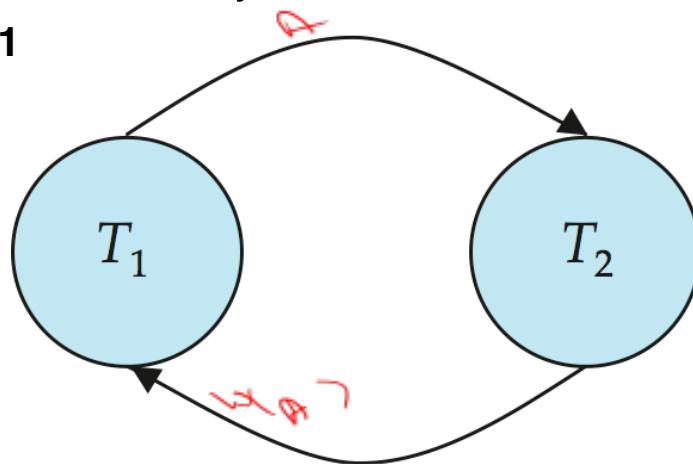
*order is according to precedence graph  
and conflicts are detected*



# Testing for Serializability

- n Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- n **Precedence graph** — a directed graph where the vertices are the transactions (names).
- n We draw an arc from  $T_i$  to  $T_j$  if the two transaction conflict, and  $T_i$  accessed the ~~data item~~ on which the conflict arose earlier.
- n We may label the arc by the item that was accessed.

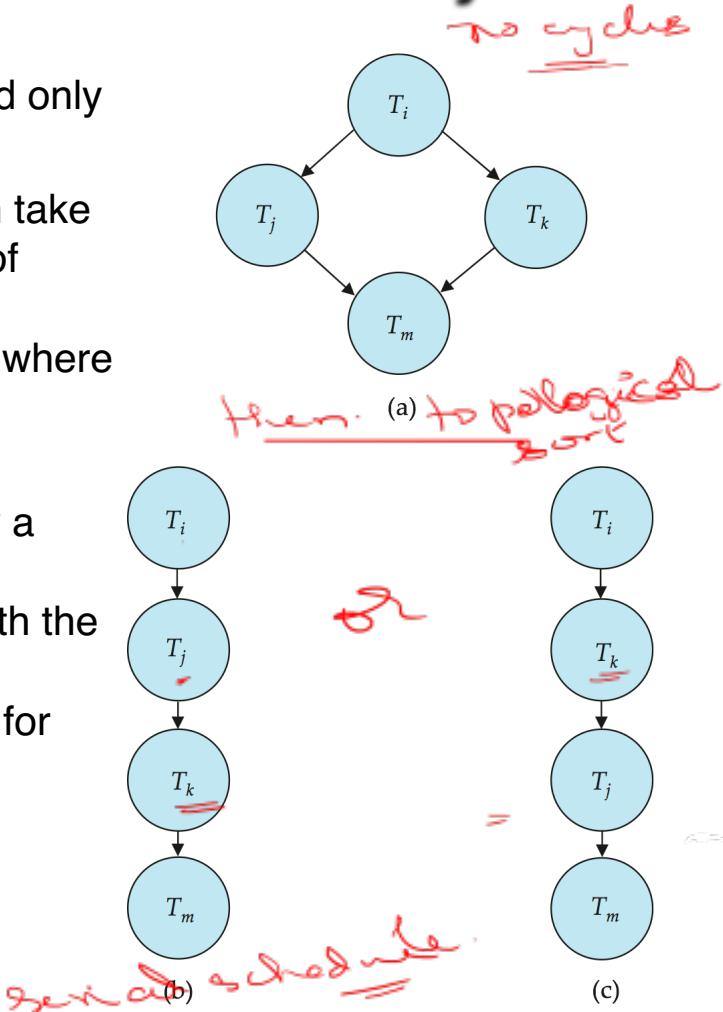
## Example 1





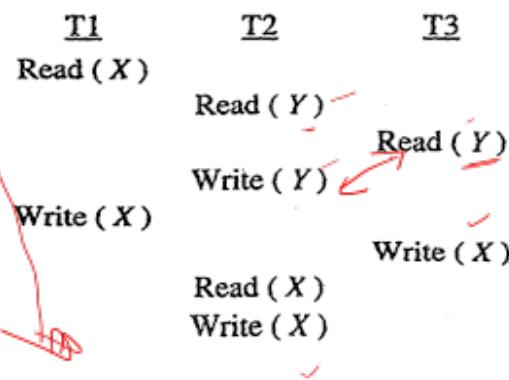
# Test for Conflict Serializability

- n A schedule is conflict serializable if and only if its precedence graph is acyclic.
- n Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph.
  - | (Better algorithms take order  $n + e$  where  $e$  is the number of edges.)
- n If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.
  - | This is a linear order consistent with the partial order of the graph.
  - | For example, a serializability order for Schedule A would be
    - T5 □ T1 □ T3 □ T2 □ T4
- 4 Are there others?

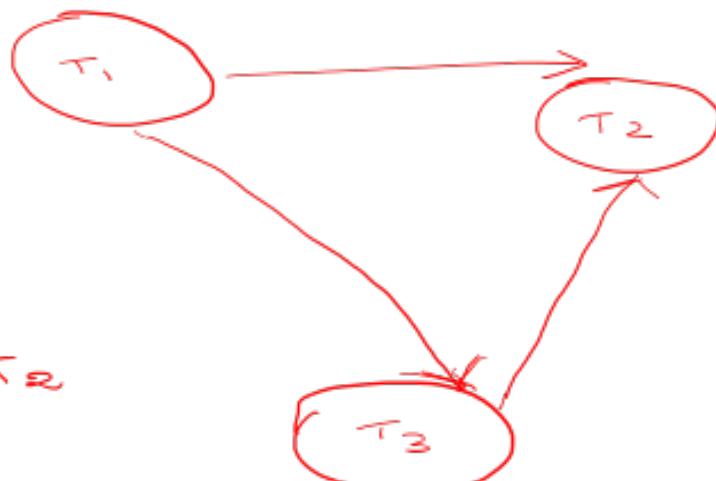




Concurrent  $S \Rightarrow$  serial schedule.



No cycles  $\therefore$  conflict  
serial schedules



Topological sort  $T_1 T_3 T_2$

serial schedule

$R(X) W(X)$ ,  $R(Y) W(Y)$ ,  $R(Y) W(Y) R(X) W(X)$

$T_1 \rightarrow T_3 \rightarrow T_2$

\* \*





Without is S<sub>1</sub> conflict serializable  
vs S<sub>2</sub> conflict non-serializable

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
w <sub>1</sub> (A)		
	R <sub>2</sub> (A)	
		w <sub>3</sub> (A)

S<sub>1</sub>

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
	w <sub>2</sub> (A)	
		w <sub>3</sub> (A)
	R <sub>2</sub> (A)	

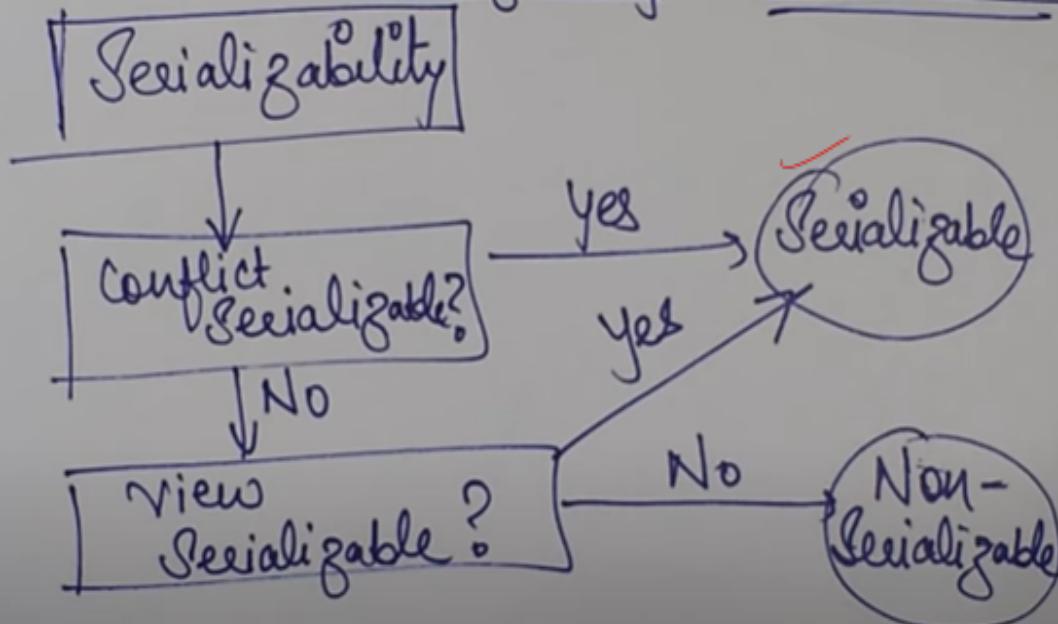
S<sub>2</sub>

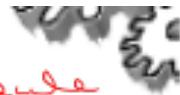
S<sub>1</sub> is not - serializable

S<sub>2</sub> not serializable =



## Serializability : A Broader View





# View Serializability

~~S<sub>1</sub> is concurrent schedule, S<sub>2</sub> is your serial schedule~~

- ♦ Schedules S<sub>1</sub> and S<sub>2</sub> are view equivalent if:

- If T<sub>i</sub> reads initial value of A in S<sub>1</sub>, then T<sub>i</sub> also reads initial value of A in S<sub>2</sub>
- If T<sub>i</sub> reads value of A written by T<sub>j</sub> in S<sub>1</sub>, then T<sub>i</sub> also reads value of A written by T<sub>j</sub> in S<sub>2</sub>
- If T<sub>i</sub> writes final value of A in S<sub>1</sub>, then T<sub>i</sub> also writes final value of A in S<sub>2</sub>

T1: R(A)	W(A)
T2:	W(A)
T3:	W(A)

S1

T1: R(A), W(A)	
T2:	W(A)
T3:	W(A)

S2





## View Serializability (Cont.)

- n A schedule  $S$  is **view serializable** if it is view equivalent to a serial schedule.
- n Every conflict serializable schedule is also view serializable.
- n Below is a schedule which is view-serializable but *not* conflict serializable.

$T_{27}$	$\cancel{T_{28}}$	$T_{29}$
read ( $Q$ )	<del>write (<math>Q</math>)</del>	
write ( $Q$ )		<del>write (<math>Q</math>)</del>

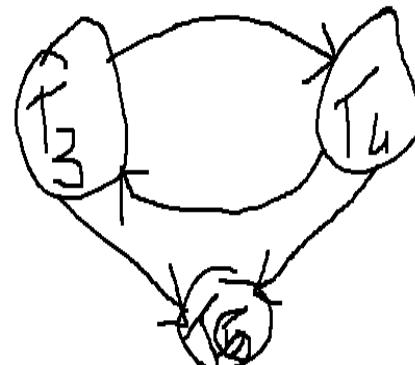
- n Every view serializable schedule that is *not* conflict serializable has **blind writes**.

without reading the item,  
the write takes place



# ex1

$T_3$	$T_4$	$T_6$
read( $Q$ )		
	write( $Q$ )	write( $Q$ )



Cycles  
not Conf/Conf

Is it view Serializable?

1. Initial read
  2. Final update
  3. writes
- $T_3 \rightarrow T_4 \rightarrow T_6$

$$T_3 \rightarrow T_4 \rightarrow T_6$$





Question :- Can this schedule be converted to a serial schedule? If so provide one.

$T_1$     $T_2$     $T_3$   
 $R(A)$

Step 1: Test 1 C.S.S. If yes then convert to a serial schedule

$w(A)$

If not step 2

$w(A)$

$w(A)$   
=

Step 2: Test view serial schedule

If not possible then no serial schedule possible



Cycles are present  
∴ no C.S.S. possible

Step 2:-

Initial Read

A

$T_1$

Find update

$T_3$

route

$T_2, T_3$

$T_1 \rightarrow T_2 \rightarrow T_3$

$R_A \quad R_B \quad W_2 \quad R \quad W_3 \quad A$

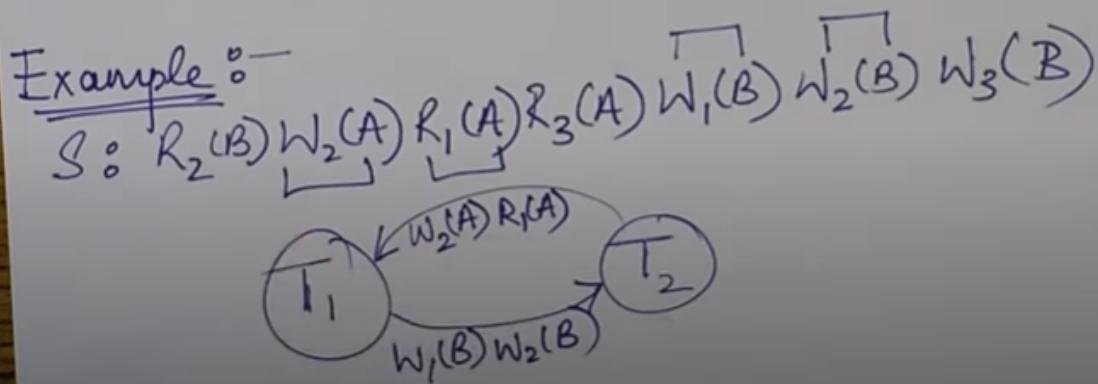
$\overbrace{T_1}^R \quad \overbrace{T_2}^W \quad \overbrace{T_3}^A \quad //$



## ex2

Schedule  $S$ , then  $T_1$  should only update finally the data item  $X$  in the view equivalent schedule  $S'$ .

Example :-





→ Non-Conflict Setting

	A	B	
Initial Read	X	T <sub>2</sub>	<del>T<sub>1</sub></del>
Final update	T <sub>2</sub>	T <sub>3</sub>	
Update	T <sub>2</sub>	T <sub>1, T<sub>2</sub>, T<sub>3</sub></sub>	$T_2 \rightarrow T_3$ $T_2 \rightarrow T_1 \rightarrow T_3$



	A	B	C
Initial Read	X	T <sub>2</sub>	T <sub>2</sub> → T <sub>3</sub>
Final Update	T <sub>2</sub>	T <sub>3</sub>	T <sub>2</sub> → T <sub>1</sub> → T <sub>3</sub>
Update	T <sub>2</sub>	T <sub>1</sub> , T <sub>2</sub> , T <sub>3</sub>	



Eg2: For view serializability

$T_1 \quad T_2 \quad T_3$

$R_B$

$R_A$

$R_A$

$R_B$

$W_B$

$W_B$

$W_B$

student wrt out  
step1: find c.s.s

====

not conflict serializable

step2:

Initial Read      A      B  
Final update       $\overline{T}_2$        $\overline{T}_3$   
write       $T_2$        $T_3$   
 $T_B \overline{T}_2 \overline{T}_3$

A      B  
 $\overline{T}_2$        $\overline{T}_3$   
 $T_2$        $T_3$   
 $T_B \overline{T}_2 \overline{T}_3$

Solution: A:  $T_2$  {  
B:  $T_2 \rightarrow T_1 \rightarrow T_3$

View serial schedule  $T_2 \rightarrow \overline{T}_1 \rightarrow \overline{T}_3$

$R_2(B) W_2(A) W_2(B) R_1(A) W_1(B) R_3(A) > W_3(B)$

=====



Solve:



Not Conflict serializable because  
of the cycles  
seen in precedence cycle

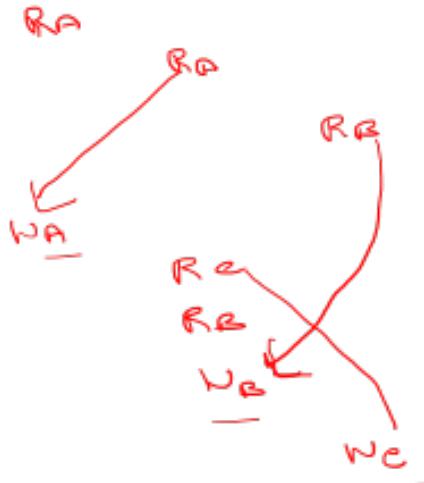
	A	B
Initial Read	$T_1, T_2$	$T_2, T_1$
Final update	$T_1$	$T_1$
writes	$T_2, T_1$	$T_1$

A:  $T_2 \rightarrow T_1$   
B:  $T_2 \rightarrow T_1$  }  $\therefore$  final serial schedule

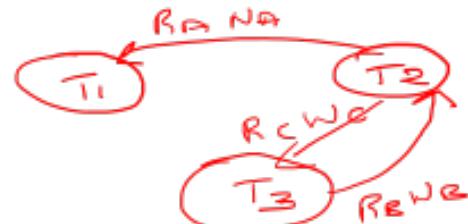
$T_2 \rightarrow T_1$   
 ~~$\leftrightarrow$~~



$T_1 \quad T_2 \quad T_3$



NOT Conflict serializable?



due to presence of cycle

Step 2:

	A	B	C
Initial Read	$T_1 \quad T_2$	$T_3 \quad T_2$	$T_2$
Final update writers	$T_1$ $T_1$	$T_2$ $T_2$	$T_3$ $T_2$

A :  $T_2 \rightarrow T_1$

B :  $T_3 \rightarrow T_2$

C :  $T_2 \rightarrow T_3$

Conflicting ~~serial~~ schedule wst B & C

∴ not view serial schedule  
possible





# ex3

Example:

T1 T2

R(A)

R(A)

W(A)

R(B)

W(A)

R(B)

W(B)

Set of all Schedules

V.S.S

C.S.  
S

Check for View serializable with  $\langle T1, T2 \rangle$  and  $\langle T2, T1 \rangle$



$T_1$        $T_2$   
R(A)  
R(A)  
W(A)  
R(B)

1. Conflict Solvable



W(A)  
R(B)  
W(B)

2. View Serializable

Initial read | A | B  
Final update |  $T_1 T_2$  |  $T_2 T_1$   
write |  $T_1$  |  $T_1$   
|  $T_2 T_1$  |  $T_1$

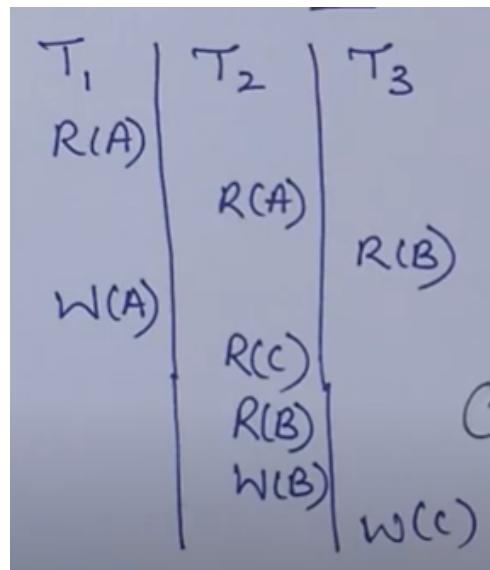
A:  $T_2 T_1$   
B:  $T_2 T_1$

$T_2 \rightarrow T_1$

	A	B
Initial read	$T_1 T_2$	$T_2 T_1$
Final update	$T_1$	$T_1$
write	$T_2 T_1$	$T_1$

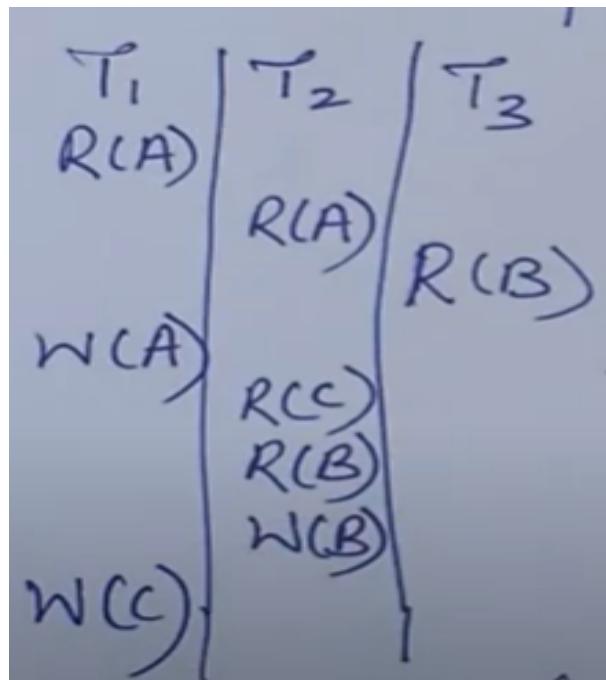


## Ex4: no serializable schedule possible





## Ex5: Serializable schedule possible





# For student workout

- n Example: Check for View equivalence

T1

R(A)

W(A)

T2

R(A)

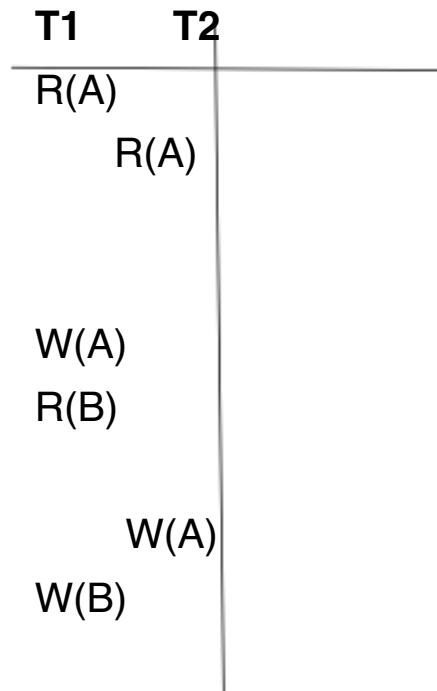
W(A)

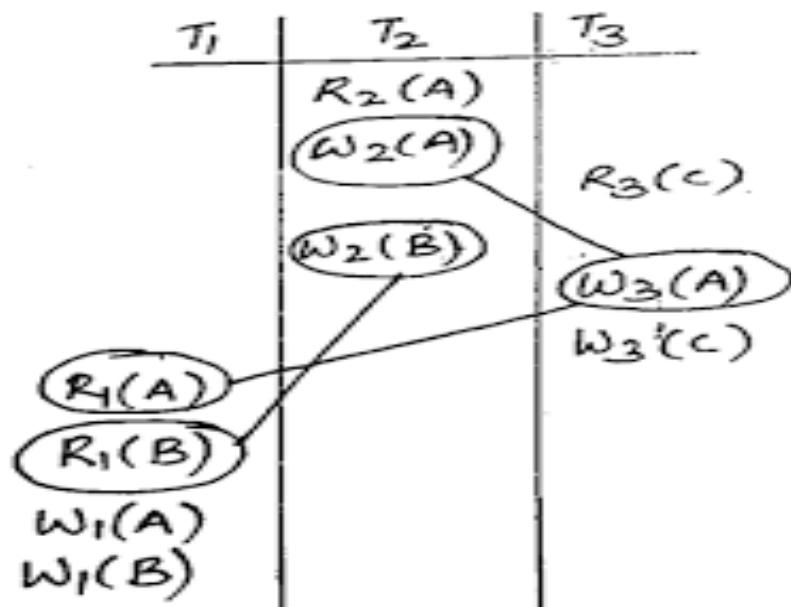
R(B)

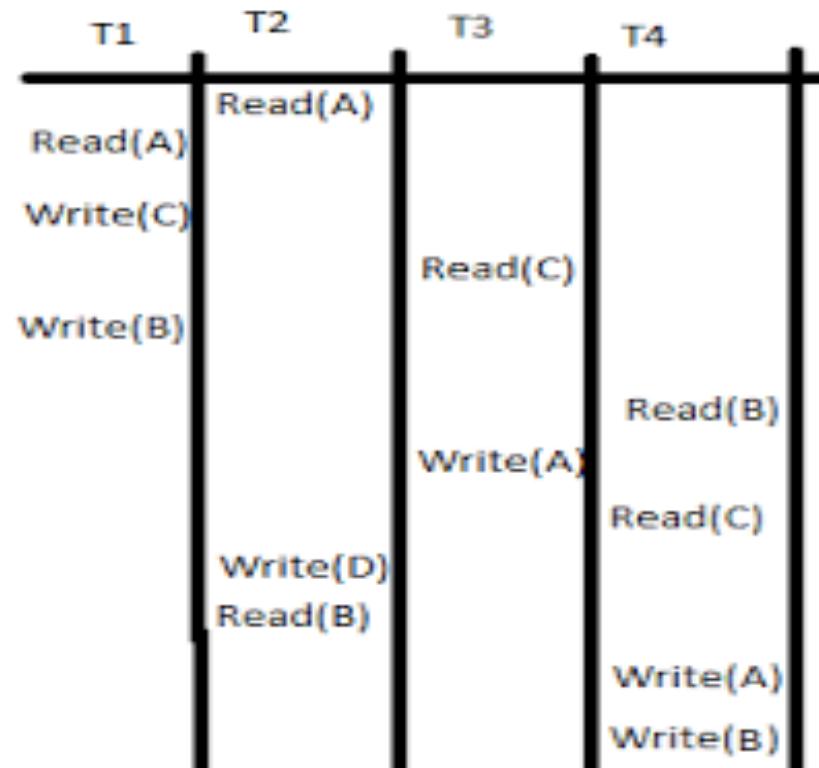
W(B)

R(B)

W(B)









# Recoverable Schedules

$T_8$	$T_9$
read ( $A$ )	
write ( $A$ )	
	read ( $A$ )
	write ( $A$ )
	commit
read ( $B$ )	
abort	

Schedule is not recoverable





# Recoverable Schedules

Need to address the effect of transaction failures on concurrently running transactions.

- n **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .
- n The following schedule (Schedule 11) is not recoverable if  $T_9$  commits immediately after the read

$T_8$	$T_9$
read ( $A$ ) write ( $A$ )	
read ( $B$ )	read ( $A$ ) commit

- n If  $T_8$  should abort,  $T_9$  would have read (and possibly shown to the user) an inconsistent database state. Hence, database must ensure that schedules are recoverable.



# Cascading Rollbacks

- n **Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks. Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read ( $A$ ) read ( $B$ ) write ( $A$ )  abort	read ( $A$ ) write ( $A$ )	read ( $A$ )

- n If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back.  
n Can lead to the **undoing of a significant amount of work**





# Cascadeless Schedules

- n **Cascade-less schedules** — cascading rollbacks cannot occur; for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$ .
- n Every cascade-less schedule is also recoverable
- n It is desirable to restrict the schedules to those that are cascade-less



# Concurrency Control

- n A database must provide a mechanism that will ensure that all possible schedules are
  - | either conflict or view serializable, and
  - | are recoverable and preferably cascadeless
- n A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
- n Testing a schedule for serializability *after* it has executed is a little too late!
- n **Goal** – to develop concurrency control protocols that will assure serializability.



# Concurrency Control (Cont.)

- n Schedules must be conflict or view serializable, and recoverable, for the sake of database consistency, and preferably cascadeless.
- n A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency.
- n Concurrency-control schemes **tradeoff between the amount of concurrency they allow and the amount of overhead that they incur.**
- n Some schemes allow only conflict-serializable schedules to be generated, while others allow view-serializable schedules that are not conflict-serializable.





# Concurrency Control vs. Serializability Tests

- n Concurrency-control protocols allow concurrent schedules, but ensure that the schedules are conflict/view serializable, and are recoverable and cascadeless .
- n Concurrency control protocols generally do not examine the precedence graph as it is being created
  - | Instead a protocol imposes a discipline that avoids nonserializable schedules.
  - | We study such protocols in Chapter 16.
- n Different concurrency control protocols provide different tradeoffs between the amount of concurrency they allow and the amount of overhead that they incur.
- n Tests for serializability help us understand why a concurrency control protocol is correct.





# Weak Levels of Consistency

- n Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
  - | E.g. a read-only transaction that wants to get an approximate total balance of all accounts
  - | E.g. database statistics computed for query optimization can be approximate (why?)
  - | Such transactions need not be serializable with respect to other transactions
- n Tradeoff accuracy for performance



# Levels of Consistency in SQL-92

- n **Serializable** — default
- n **Repeatable read** — only committed records to be read, repeated reads of same record must return same value. However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others.
- n **Read committed** — only committed records can be read, but successive reads of record may return different (but committed) values.
- n **Read uncommitted** — even uncommitted records may be read.
  
- n Lower degrees of consistency useful for gathering approximate information about the database
- n Warning: some database systems do not ensure serializable schedules by default
  - | E.g. Oracle and PostgreSQL by default support a level of consistency called snapshot isolation (not part of the SQL standard)



# Transaction Definition in SQL

- n Data manipulation language must include a construct for specifying the set of actions that comprise a transaction.
- n In SQL, a transaction begins implicitly.
- n A transaction in SQL ends by:
  - | **Commit work** commits current transaction and begins a new one.
  - | **Rollback work** causes current transaction to abort.
- n In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
  - | Implicit commit can be turned off by a database directive
    - 4 E.g. in JDBC, connection.setAutoCommit(false);





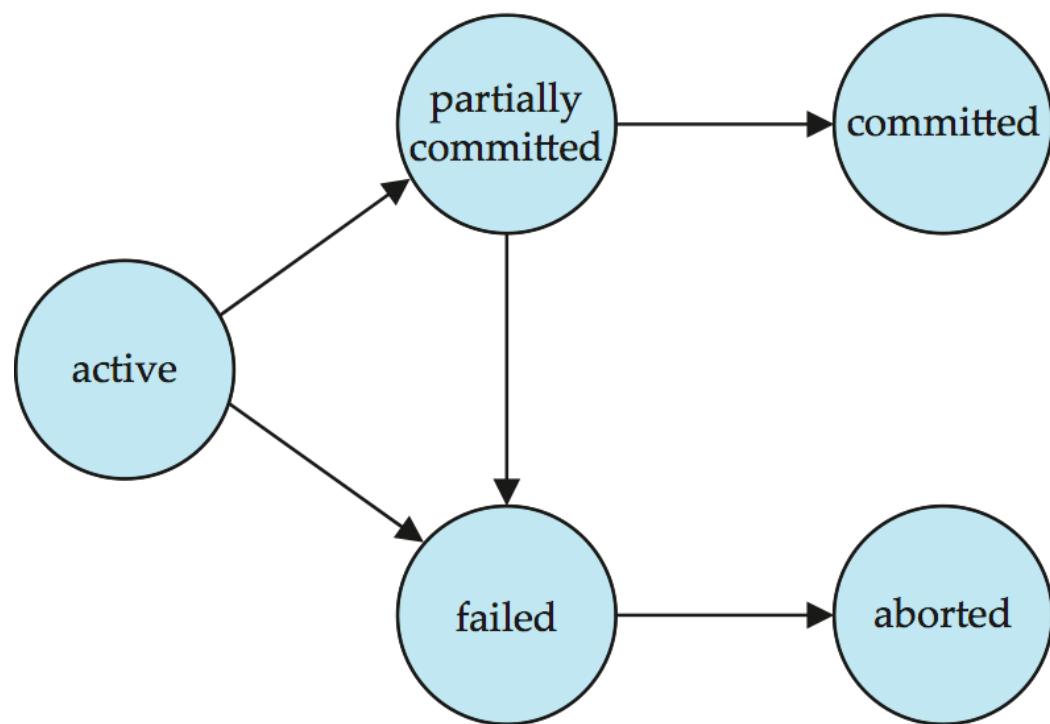
## **End of Chapter 14**

**Database System Concepts, 6th  
Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Figure 14.01





# Figure 14.02

$T_1$	$T_2$
<p>read (<math>A</math>) <math>A := A - 50</math> write (<math>A</math>) read (<math>B</math>) <math>B := B + 50</math> write (<math>B</math>) commit</p>	<p>read (<math>A</math>) <math>temp := A * 0.1</math> <math>A := A - temp</math> write (<math>A</math>) read (<math>B</math>) <math>B := B + temp</math> write (<math>B</math>) commit</p>



# Figure 14.03

$T_1$	$T_2$
	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	



# Figure 14.04

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ )  read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ )  read ( $B$ ) $B := B + temp$ write ( $B$ ) commit



# Figure 14.05

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$  write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ )  $B := B + temp$ write ( $B$ ) commit



# Figure 14.06

$T_1$	$T_2$
read ( $A$ ) write ( $A$ )	read ( $A$ ) write ( $A$ )
read ( $B$ ) write ( $B$ )	read ( $B$ ) write ( $B$ )



# Figure 14.07

$T_1$	$T_2$
read ( $A$ )	
write ( $A$ )	
read ( $B$ )	read ( $A$ )
write ( $B$ )	write ( $A$ )
	read ( $B$ )
	write ( $B$ )





# Figure 14.08

$T_1$	$T_2$
read ( $A$ )	
write ( $A$ )	
read ( $B$ )	
write ( $B$ )	read ( $A$ )
	write ( $A$ )
	read ( $B$ )
	write ( $B$ )

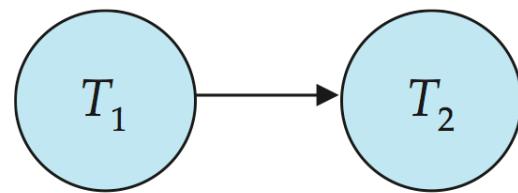


# Figure 14.09

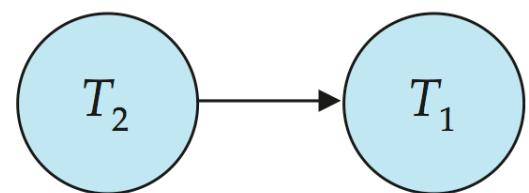
$T_3$	$T_4$
read ( $Q$ )	
write ( $Q$ )	write ( $Q$ )



# Figure 14.10



(a)

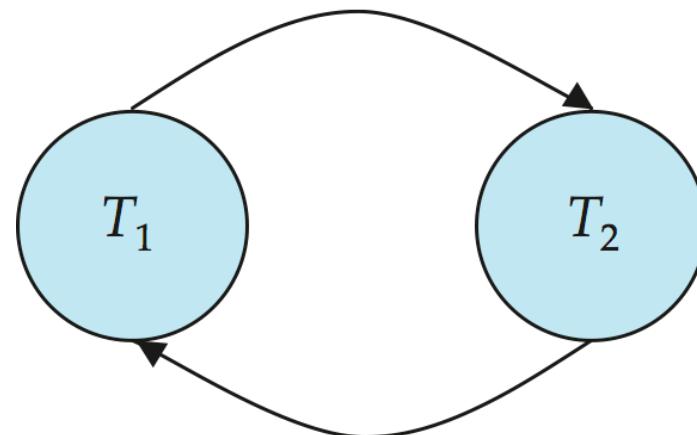


(b)



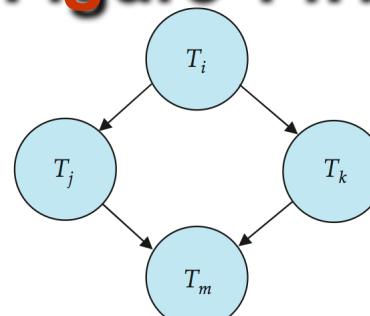


# Figure 14.11

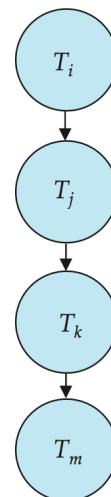




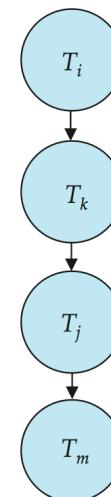
# Figure 14.12



(a)



(b)



(c)



# Figure 14.13

$T_1$	$T_5$
read ( $A$ ) $A := A - 50$ write ( $A$ )	
read ( $B$ ) $B := B + 50$ write ( $B$ )	read ( $B$ ) $B := B - 10$ write ( $B$ )
	read ( $A$ ) $A := A + 10$ write ( $A$ )



# Figure 14.14

$T_8$	$T_9$
read ( $A$ ) write ( $A$ )	
read ( $B$ )	read ( $A$ ) commit



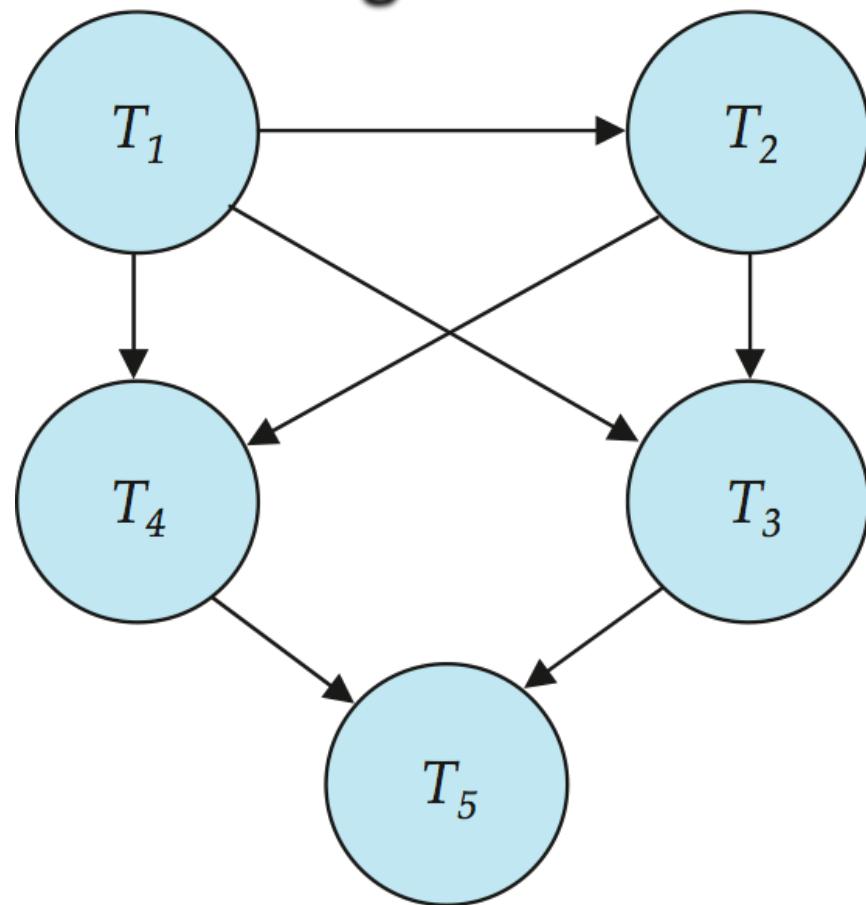
## Figure 14.15

$T_{10}$	$T_{11}$	$T_{12}$
read ( $A$ ) read ( $B$ ) write ( $A$ )  abort	read ( $A$ ) write ( $A$ )	read ( $A$ )





# Figure 14.16





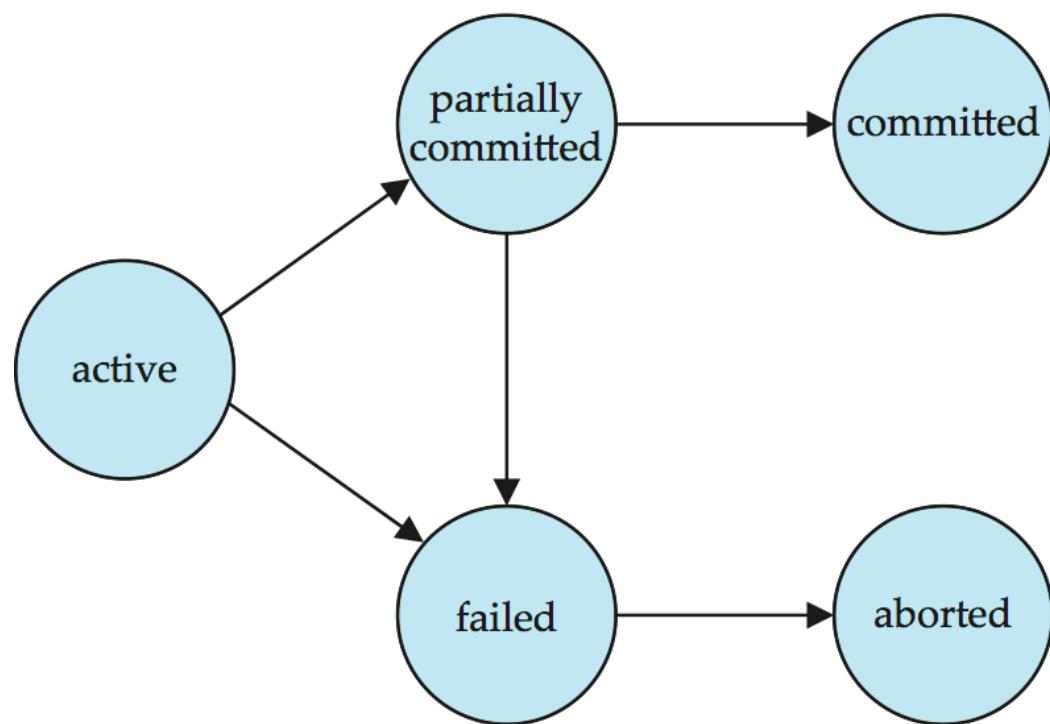
# **End of Chapter 14**

**Database System Concepts, 6th  
Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Figure 14.01





# Figure 14.02

$T_1$	$T_2$
<p>read (<math>A</math>) <math>A := A - 50</math> write (<math>A</math>) read (<math>B</math>) <math>B := B + 50</math> write (<math>B</math>) commit</p>	<p>read (<math>A</math>) <math>temp := A * 0.1</math> <math>A := A - temp</math> write (<math>A</math>) read (<math>B</math>) <math>B := B + temp</math> write (<math>B</math>) commit</p>



# Figure 14.03

$T_1$	$T_2$
	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	



# Figure 14.04

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ )  read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ )  read ( $B$ ) $B := B + temp$ write ( $B$ ) commit



# Figure 14.05

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$  write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ )  $B := B + temp$ write ( $B$ ) commit



# Figure 14.06

$T_1$	$T_2$
read ( $A$ ) write ( $A$ )	read ( $A$ ) write ( $A$ )
read ( $B$ ) write ( $B$ )	read ( $B$ ) write ( $B$ )





# Figure 14.07

$T_1$	$T_2$
read ( $A$ )	
write ( $A$ )	
read ( $B$ )	read ( $A$ )
write ( $B$ )	write ( $A$ )
	read ( $B$ )
	write ( $B$ )





# Figure 14.08

$T_1$	$T_2$
read ( $A$ )	
write ( $A$ )	
read ( $B$ )	
write ( $B$ )	read ( $A$ )
	write ( $A$ )
	read ( $B$ )
	write ( $B$ )

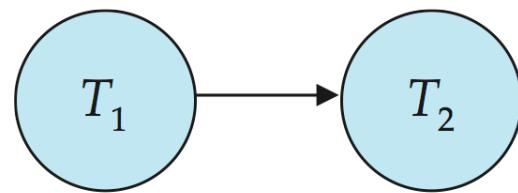


# Figure 14.09

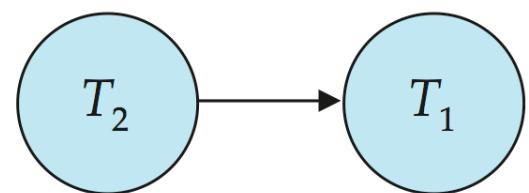
$T_3$	$T_4$
read ( $Q$ )	
write ( $Q$ )	write ( $Q$ )



# Figure 14.10



(a)

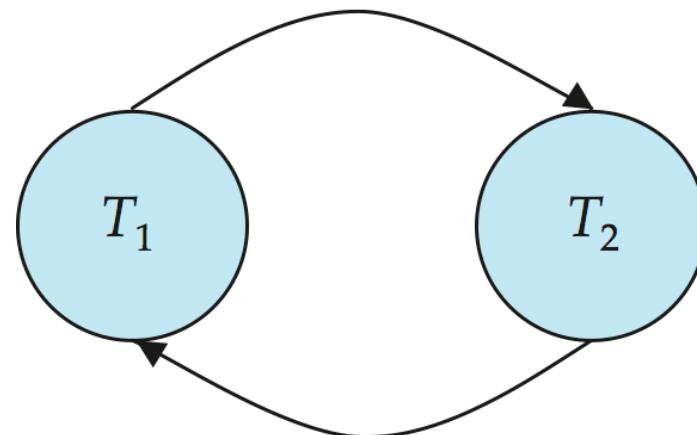


(b)



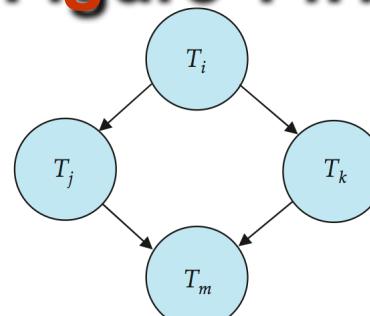


# Figure 14.11

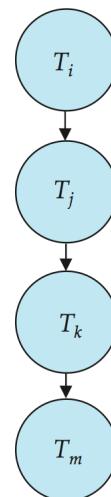




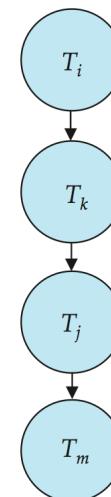
# Figure 14.12



(a)



(b)



(c)



# Figure 14.13

$T_1$	$T_5$
read ( $A$ ) $A := A - 50$ write ( $A$ )	
read ( $B$ ) $B := B + 50$ write ( $B$ )	read ( $B$ ) $B := B - 10$ write ( $B$ )
	read ( $A$ ) $A := A + 10$ write ( $A$ )



# Figure 14.14

$T_8$	$T_9$
read ( $A$ ) write ( $A$ )	
read ( $B$ )	read ( $A$ ) commit



## Figure 14.15

$T_{10}$	$T_{11}$	$T_{12}$
read ( $A$ ) read ( $B$ ) write ( $A$ )  abort	read ( $A$ ) write ( $A$ )	read ( $A$ )





# Figure 14.16

