

Basic sql- part 3

Natural join to null values in aggregate function

Natural Join

- ❑ Natural join matches tuples with the same values for all common attributes, and retains only one copy of each common column
- ❑ **select * from instructor natural join teaches;**

instructor *teaches*

ID	name	dept_name	salary	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Finance	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	El Said	History	60000	HIS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010
83821	Brandt	Comp. Sci.	92000	CS-190	1	Spring	2009
83821	Brandt	Comp. Sci.	92000	CS-190	2	Spring	2009
83821	Brandt	Comp. Sci.	92000	CS-319	2	Spring	2010
98345	Kim	Elec. Eng.	80000	EE-181	1	Spring	2009

Natural Join Example

- ② List the names of instructors along with the course ID of the courses that they taught.
- ② In two ways we can write the Query:

② **select name, course_id from instructor, teaches
where instructor.ID = teaches.ID;**



OR



② **select name, course_id
from instructor natural join teaches;**

- A **from** clause in an SQL query can have multiple relations combined using natural join, as shown here:
SQL
- **select** A_1, A_2, \dots, A_n
- **from** r_1 natural join r_2 natural join \dots natural join r_m
- **where** $P;$

- Retrieve the names of all instructors, along with their department names and department building name.
- instructor(ID, name, dept_name, salary)
- department(dept_name, building, budget)

- Query:
select name, dept_name, building from instructor natural join department;
- Select name, dept_name , building from instructor, department where instructor.dept_name=department.dept_name
- Using natural join

- Display ~~the~~ **computer science courses** which are held in **room no=105** and building ~~AB5~~
- course(course_id, title, dept_name, credits)
- section(course_id, sec_id, semester, year, building, room_number, time_slot_id)
- select ~~course_id~~ from course natural join section
course.course_id=section.course_id and building='AB5' and
room_number=105 and dept_name='CS';
using natural join:

Natural Join (Cont.)

- ❑ Danger in natural join: beware of unrelated attributes with same name which get equated incorrectly
- ❑ ~~course(course_id, title, dept_name, credits)~~
- ❑ ~~teaches(ID, course_id, sec_id, semester, year)~~
- ❑ ~~instructor(ID, name, dept_name, salary)~~
- ❑ List the names of instructors along with the titles of courses that they teach
 - ❑ **select name, title
from instructor natural join teaches natural join course;**

the natural join of instructor and teaches contains the attributes (ID, name, dept name, salary, course id, sec id, sem, year), while the course relation contains the attributes (course id, title, dept name, credits). As a result, the natural join of these two would require that the dept name attribute values from the two inputs be the same, in addition to requiring that the course id values be the same. This query would then omit all (instructor name, course title) pairs where the instructor teaches a course in a department other than the instructor's own department.

Instructor natural join teaches

- (ID, name, dept name, salary, course id, sec id, sem, year), 

Course relations is as follows :

- (course id, title, dept name, credits). 

If  natural joins 

common attributes courseid, deptname



- This query would then omit all (instructor name, course title) pairs where the instructor teaches a course in a department other than the instructor's own department.

~~instructor . dept - name = course . dept - name~~ X

List the names of instructors along with the titles of courses that they teach



correct version

select name, title
from (instructor natural join teaches)
join course using(course_id);

list of attr.
X



- The operation join . . . using requires a list of attribute names to be specified.
one : R₁ C₄ - R₂ C₅
- No other common attribute even if present will be matched, other than the ones indicated in using clause
select A from R₁ join R₂ using B₃

- Display *id* of students who have taken the course offered by physics department
- student(SID, name, dept_name, tot_cred)
- takes(SID, course_id, sec_id, semester, year, grade)
- course(course_id, title, dept_name, credits)
-
-
- Select student.SID from student, course, takes where student.SID=takes.SID and course.course_id =takes.course_id and course.dept_name="physics" ;
- **Select student.SID from (student natural join takes),join course using course_id where course.dept_name="physics"**

- Display the student' details who scored A or A+ grade in "DBS" course.
- student(SID, name, dept_name, tot_cred)
- takes(SID, course_id, sec_id, semester, year, grade)
~~student * sec from student natural joins takes~~
- course(course_id, title, dept_name, credits)
~~join course using (course_id)~~

where grade = "A" or grade = "A+"
and title = "DBS"

Additional Basic Operations

The Rename Operation

The names of the attributes in the result are derived from the names of the attributes in the relations in the **from** clause.

We may need to change this for the following reasons:

- two relations in the **from** clause may have attributes with the same name, in which case an attribute name is duplicated in the result.
 - if we used an arithmetic expression in the **select** clause, the resultant attribute does not have a name.
 - even if an attribute name can be derived from the base relations as in the preceding example, we may want to change the attribute name in the result.
- ?
- The SQL allows renaming relations and attributes using the **as** clause:

old-name as new-name

?

E.g.

?

select ID, name, salary/12 as monthly_salary from instructor

ID	name	salary / 12

After rename

ID	name	<u>monthlysal</u>

- One reason to rename a relation is to replace a long relation name with a shortened version that is more convenient to use elsewhere in the query.
- **select** *T.name, s.course id*
- **from** *instructor as T, teaches as s*
- **where** *T.ID= s.ID;*

Another reason to rename a relation is a case where we wish to compare tuples in the same relation. We then need to take the Cartesian product of a relation with itself and, without renaming, it becomes impossible to distinguish one tuple from the other.

- ② Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

ID	name	dept	salary
x			X
x			X
x			X
x			X

ID	name	dept	salary
a	a	a	a
b	b	b	b
c	c	c	c
d	d	d	d

② select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'

An identifier, such as T and S , that is used to rename a relation is referred to as a correlation name in the SQL standard, but is also commonly referred to as a table alias, or a correlation variable, or a tuple variable.

String Operations

Note: The SQL standard specifies that the equality operation on strings is case sensitive; as a result the expression “comp. sci.” = ‘Comp. Sci.’” **evaluates to false**. However, some database systems, such as **MySQL** and **SQL Server**, do not distinguish uppercase from lowercase when matching strings; as a result “comp. sci.”= ‘Comp. Sci.’” would evaluate to **true** on these databases. This

- ❑ SQL includes a string-matching operator for comparisons on character strings. The operator “**like**” uses patterns that are described using two special characters:
 - ❑ percent (%). The % character matches **any substring**.
 - ❑ underscore (_). The _ character matches **any character**.
- ❑ Find the names of all instructors whose **name includes the substring “dar”**.
select name from instructor where name like '%dar%'

- ? Patters are **case sensitive**.
- ? Pattern matching examples:
 - ? ‘Intro%’ matches any string **beginning with “Intro”**.
 - ? ‘%Comp%’ matches any string **containing “Comp” as a substring**.
 - ? ‘___’ matches any string of **exactly three characters**.
 - ? ‘___ %’ matches any string of at **least three characters**.

- display all the records from the CUSTOMERS table, where
- exercises** the SALARY starts with 200.

```
SELECT * FROM table_name WHERE column LIKE 'X%'  
SELECT * FROM table_name WHERE column LIKE '%X%'  
SELECT * FROM table_name WHERE column LIKE 'X_'  
SELECT * FROM table_name WHERE column LIKE '_X'  
SELECT * FROM table_name WHERE column LIKE '_X_'
```

String Operations (Cont.)

- ② Match the string “100 %” // Usage of % in a string
- ② ‘\’ is used to specify the special character such as % or \\.
(eg. Retrieve whose attendance status is 100%)
Select names from student where attendance like ‘100%’

“\” implies to different things to different

- ② SQL supports a variety of **string operations** such as
 - ② concatenation (using “||”)
 - ② converting from **upper to lower case** (and vice versa)
 - ② finding string **length**, **extracting substrings**, etc.

Eg.: Select **SUBSTR**('ABCDEF', 2, 3) as substring from tablename ; o/p: BCD

```
select length(name) as length_val from dual ;
```

```
select upper('hi') as Upper_val from dual ;
```

Ordering the Display of Tuples

- >List in alphabetic order the names of all instructors

```
select distinct name  
from instructor  
order by name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

- Example: **order by name desc**

- Can sort on multiple attributes

- Example: **order by dept_name, name**

- General Form:**

- **SELECT distinct name**
 - FROM instructor**
 - **WHERE Dept_name='CS'**
 - **ORDER BY name ;**

Where Clause Predicates

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is, $\geq \$90,000$ and $\leq \$100,000$)

```
select name  
from instructor  
where salary between 90000 and 100000
```

- Tuple comparison
- ```
select name, course_id
from instructor, teaches
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```
- All SQL platforms may not support this syntax



# Set Operations

- Find courses that ran in Fall 2009 or in Spring 2010

```
(select course_id from section where sem = 'Fall' and year = 2009)
union
(select course_id from section where sem = 'Spring' and year = 2010)
```

Interpretation is important!



- Find courses that ran in Fall 2009 and in Spring 2010
- (**select course\_id from section where sem = 'Fall' and year = 2009**)  
**intersect**  
**(select course\_id from section where sem = 'Spring' and year = 2010)**

- Find courses that ran in Fall 2009 but not in Spring 2010
- (**select course\_id from section where sem = 'Fall' and year = 2009**)  
**except**
- (**select course\_id from section where sem = 'Spring' and year = 2010**)

# Set Operations

- ② Set operations **union**, **intersect**, and **except**
  - ③ Each of the above operations automatically eliminates duplicates
- ② To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

## Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary ✓

```
select distinct T.salary
from instructor as T, instructor as S
where T.salary < S.salary
```

- Find all the salaries of all instructors ✓

```
select distinct salary
from instructor
```

- Find the largest salary of all instructors

```
(select "second query")
except
(select "first query")
```

# Null Values

- ❑ It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- ❑ *null* signifies an unknown value or that a value does not exist.
- ❑ The result of any arithmetic expression involving *null* is *null*
  - ❑ Example:  $5 + \text{null}$  returns *null*
- ❑ The predicate **is null** can be used to check for null values.
  - ❑ Example: Find all instructors whose salary is null.

```
select name from instructor where salary is null;
```

| ID  | Student         | Email1                       | Email2                | Phone         | Father            |
|-----|-----------------|------------------------------|-----------------------|---------------|-------------------|
| 1   | Stan Marsh      | NULL                         | smarsh@gmail.com      | 740-097-0951  | Randy Marsh       |
| 2   | Butters Stotch  | bstotch@spelementary.com     | bstotch@gmail.com     | NULL          | Stephen Stotch    |
| 3   | Kenny McCormick | NULL                         | NULL                  |               | Stuart McCormick  |
| 4   | Kyle Broflovski | kbroflovski@spelementary.com | kbroflovski@gmail.com | 619-722-2491  | Gerald Broflovski |
| 666 | Eric Cartman    | ecartman@spelementary.com    | ecartman@gmail.com    | 740-6733-5671 | NULL              |

| ID | Student         | Email1 | Email2 |
|----|-----------------|--------|--------|
| 3  | Kenny McCormick | NULL   | NULL   |

- **SELECT ID, Student, Email1, Email2  
FROM tblSouthPark WHERE Email1 IS NULL AND Email2 IS NULL;**

# Null Values and Three Valued Logic

- ② Any comparison with *null* returns *unknown*
  - ② Example:  $5 < \text{null}$  or  $\text{null} < > \text{null}$  or  $\text{null} = \text{null}$
- ② Three-valued logic using the truth value *unknown*:
  - ② OR: (*unknown or true*) = *true*,  
          (*unknown or false*) = *unknown*  
          (*unknown or unknown*) = *unknown*
  - ② AND: (*true and unknown*) = *unknown*,  
          (*false and unknown*) = *false*,  
          (*unknown and unknown*) = *unknown*
  - ② NOT: (*not unknown*) = *unknown*
  - ② “*P is unknown*” evaluates to true if P is not known

# Aggregate Functions

*Aggregate functions* are functions that take a collection (a set or multiset) of values as input and return a single value. SQL offers five built-in aggregate functions:

**avg**: average value

**min**: minimum value

**max**: maximum value

**sum**: sum of values

**count**: number of values

# Aggregate Functions (Cont.)

- ② Find the average salary of instructors in the Computer Science department

```
② select avg (salary)
 from instructor
 where dept_name= 'Comp. Sci.';
```

- ② Find the total number of instructors who teach a course in the Spring 2010 semester

```
② select count (distinct ID) [Distinct: Since same teacher teaching more
 than one subject in Spring 10]
 from teaches
 where semester = 'Spring' and year = 2010
```

- ② Find the number of tuples in the *course* relation

```
② select count (*)
 from course;
```

# “group by” clause

- There are circumstances where we would like to apply the aggregate function not only to a single set of tuples, but also to a group of sets of tuples; we specify this wish in SQL using the **group by** clause.
- The attribute or attributes given in the **group by** clause are used to form groups.
- Tuples with the same value on all attributes in the **group by** clause are placed in one group.

# Aggregate Functions – Group By

- Find the average salary of instructors in each department

- `select dept_name, avg (salary) from instructor  
group by dept_name;`

- Note: departments with no instructor will not appear in result

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 76766 | Crick      | Biology    | 72000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 12121 | Wu         | Finance    | 90000  |
| 76543 | Singh      | Finance    | 80000  |
| 32343 | El Said    | History    | 60000  |
| 58583 | Califieri  | History    | 62000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 22222 | Einstein   | Physics    | 95000  |

| dept_name  | avg_salary |
|------------|------------|
| Biology    | 72000      |
| Comp. Sci. | 77333      |
| Elec. Eng. | 80000      |
| Finance    | 85000      |
| History    | 61000      |
| Music      | 40000      |
| Physics    | 91000      |

- Find the number of instructors in each department who teach a course in the Spring 2
- **select *dept name*, count (distinct *ID*) as *instr\_count***
- **from *instructor* natural**
- **where *semester* = 'Sprin**
- **group by *dept name*; 0**

| <i>dept_name</i> | <i>instr_count</i> |
|------------------|--------------------|
| Comp. Sci.       | 3                  |
| Finance          | 1                  |
| History          | 1                  |
| Music            | 1                  |

# Aggregation (Cont.)

- ② Attributes in **select** clause outside of aggregate functions must appear in **group by** list.

In other words, any attribute that is not present in the **group by** clause must appear only inside an aggregate function if it appears in the **select** clause, otherwise the query is treated as erroneous.

```
③ /* erroneous query */
select dept_name, ID, avg (salary)
from instructor
group by dept_name;
```

Each instructor in a particular group (defined by *dept name*) can have a different *ID*, and since only one tuple is output for each group, there is no unique way of choosing which *ID* value to output. As a result, such cases are disallowed by SQL.

# exercise

---

| ord_no | purch_amt | ord_date   | customer_id | salesman_id |
|--------|-----------|------------|-------------|-------------|
| 70001  | 150.5     | 2012-10-05 | 3005        | 5002        |
| 70009  | 270.65    | 2012-09-10 | 3001        | 5005        |
| 70002  | 65.26     | 2012-10-05 | 3002        | 5001        |
| 70004  | 110.5     | 2012-08-17 | 3009        | 5003        |
| 70007  | 948.5     | 2012-09-10 | 3005        | 5002        |
| 70005  | 2400.6    | 2012-07-27 | 3007        | 5001        |
| 70008  | 5760      | 2012-09-10 | 3002        | 5001        |

- Write a SQL statement to find the total purchase amount of all orders

Select sum(purch\_amt) from orders;

- Write a SQL statement to find the number of salesmen currently listing for all of their customers.

Select count(distinct salesman\_id) from orders;

| ord_no | purch_amt | ord_date   | customer_id | salesman_id |
|--------|-----------|------------|-------------|-------------|
| 70001  | 150.5     | 2012-10-05 | 3005        | 5002        |
| 70009  | 270.65    | 2012-09-10 | 3001        | 5005        |
| 70002  | 65.26     | 2012-10-05 | 3002        | 5001        |
| 70004  | 110.5     | 2012-08-17 | 3009        | 5003        |
| 70007  | 948.5     | 2012-09-10 | 3005        | 5002        |
| 70005  | 2400.6    | 2012-07-27 | 3007        | 5001        |
| 70008  | 5760      | 2012-09-10 | 3002        | 5001        |

```
1 | SELECT customer_id,ord_date,MAX(purch_amt)
2 | FROM orders
3 | GROUP BY customer_id,ord_date;
```

Output of the Query:

| customer_id | ord_date   | max     |
|-------------|------------|---------|
| 3002        | 2012-10-05 | 65.26   |
| 3003        | 2012-08-17 | 75.29   |
| 3005        | 2012-10-05 | 150.50  |
| 3007        | 2012-07-27 | 2400.60 |
| 3009        | 2012-08-17 | 110.50  |
| 3001        | 2012-09-10 | 270.65  |
| 3002        | 2012-09-10 | 5760.00 |
| 3005        | 2012-09-10 | 948.50  |
| 3009        | 2012-10-10 | 2480.40 |
| 3008        | 2012-06-27 | 250.45  |
| 3004        | 2012-10-10 | 1983.43 |
| 3002        | 2012-04-25 | 3045.60 |

| customer_id | cust_name      | city       | grade | salesman_id |
|-------------|----------------|------------|-------|-------------|
| 3002        | Nick Rimando   | New York   | 100   | 5001        |
| 3007        | Brad Davis     | New York   | 200   | 5001        |
| 3005        | Graham Zusi    | California | 200   | 5002        |
| 3008        | Julian Green   | London     | 300   | 5002        |
| 3004        | Fabian Johnson | Paris      | 300   | 5006        |
| 3009        | Geoff Cameron  | Berlin     | 100   | 5003        |
| 3003        | Jozv Altidor   | Moscow     | 200   | 5007        |

- Write a SQL statement to know how many customer have listed their names.
- find the number of customers who gets at least a gradation for his/her performance.
- find the highest purchase amount by the each customer with their ID and highest purchase amount
- `SELECT customer_id, MAX(purch_amt) FROM orders GROUP BY customer_id;`

# “Group by... having “ clause

- At times, it is useful to state a condition that applies to groups rather than to tuples.
- For example, we might be interested in only those departments where the average salary of the instructors is more than \$42,000.
- This condition does not apply to a single tuple; rather, it applies to each group constructed by the **group by** clause.
- To express such a query, we use the **having** clause of SQL. SQL applies predicates in the **having** clause after groups have been formed, so aggregate functions may be used

# Aggregate Functions – Having Clause

- ② Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

| dept_name  | avg(avg_salary) |
|------------|-----------------|
| Physics    | 91000           |
| Elec. Eng. | 80000           |
| Finance    | 85000           |
| Comp. Sci. | 77333           |
| Biology    | 72000           |
| History    | 61000           |

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

```
SELECT commission, COUNT (*)
FROM agents
GROUP BY commission
HAVING COUNT (*) > 3;
```

agents

| AGENT_NAME | COMMISSION |
|------------|------------|
| Alex       | .13        |
| Subbarao   | .14        |
| Benjamin   | .11        |
| Ramasundar | .15        |
| Alford     | .12        |
| Ravi Kumar | .15        |
| Santakumar | .14        |
| Lucida     | .12        |
| Anderson   | .13        |
| Mukesh     | .11        |
| McDen      | .15        |
| Ivan       | .15        |

GROUP BY commission

| COMMISSION | COUNT(*) |
|------------|----------|
| .15        | 4        |
| .11        | 2        |
| .14        | 2        |
| .13        | 2        |
| .12        | 2        |

HAVING COUNT (\*) > 3;

| COMMISSION | COUNT(*) |
|------------|----------|
| .15        | 4        |

- To illustrate the use of both a **having** clause and a **where** clause in the same query, we consider the query “For each course section offered in 2009, find the average total credits (*tot cred*) of all students enrolled in the section, if the section had at least 2 students.”
- **select** *course\_id, semester, year, sec\_id, avg (tot cred)*
- **from** *takes natural join student*
- **where** *year = 2009*
- **group by** *course id, semester, year, sec\_id*
- **having count (ID) >= 2;**

# Null Values and Aggregates

- ❑ Total all salaries

```
select sum (salary)
from instructor
```

- ❑ Above statement ignores null amounts
- ❑ All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- ❑ What if collection has only null values?
  - ❑ count returns 0
  - ❑ all other aggregates return null
- ❑

**Table1:**

| Field1 | Field2 | Field3 |
|--------|--------|--------|
| 1      | 1      | 1      |
| NULL   | NULL   | NULL   |
| 2      | 2      | NULL   |
| 1      | 3      | 1      |

Then

```
SELECT COUNT(*), COUNT(Field1), COUNT(Field2), COUNT(DISTINCT Field3)
FROM Table1
```

Output Is:

```
COUNT(*) = 4; -- count all rows even null/duplicates

-- count only rows without null values on that field
COUNT(Field1) = COUNT(Field2) = 3

COUNT(Field3) = 2
COUNT(DISTINCT Field3) = 1 -- Ignore duplicates
```