



# Chapter 4: Intermediate SQL

## part1- join to views

# Chapter 4: Intermediate SQL

Join Expressions

- Views
- Transactions
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization

# Joined Relations

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause

```
select name, course_id  
from instructor natural join teaches;
```

# Joined Relations

- Join operations take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

n Suppose we wish to display a list of all students, displaying their *ID*, and *name*, *dept name*, and *tot cred*, along with the courses that they have taken. The following SQL query may appear to retrieve the required information:

```
select * from student natural join takes;
```

We would thus not see any information about students who have not taken any course.

# Join operations – Example

Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

n Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

n Observe that

prereq information is missing for CS-315 and  
course information is missing for CS-437

# Natural or conditional Joins

*n course natural join prereq*

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101

# Outer Join

- More generally, some tuples in either or both of the relations being joined may be “lost” in this way.
- An extension of the join operation that **avoids loss of information**.
- The **outer join** operation works in a manner similar to the join operations we have already studied, but preserve those tuples that would be lost in a join, by creating **tuples in the result containing null values**.
- Computes the join and then adds tuples from one relation that **does not match tuples** in the other relation to the result of the join.
- Uses *null* values.

- There are in fact three forms of outer join:
- The **left outer join** preserves tuples only in the relation named before (to the left of) the **left outer join** operation.
- The **right outer join** preserves tuples only in the relation named after (to the right of) the **right outer join** operation.
- The **full outer join** preserves tuples in both relations.

Join types	Join Conditions
inner join	natural <del>common</del>
left outer join	on <u>&lt; predicate &gt;</u>
right outer join	using <u>(A<sub>1</sub>, A<sub>1</sub>, ..., A<sub>n</sub>)</u>
full outer join	<del>left join</del>

# Left Outer Join

*n course natural left outer join prereq*

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null

# Right Outer Join

n course **natural right outer join**  
*prereq*

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

# Full Outer Join

n course **natural full outer join**  
prereq

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101

- “Display a list of all students in the Comp. Sci. department, along with the course sections, if any, that they have taken in Spring 2009; all course sections from Spring 2009 must be displayed, even if no student from the Comp. Sci. department has taken the course section.” This query can be written as:
- **select \* from (select \* from student where dept name= 'Comp. Sci')**  
**natural full outer join (select \* from takes where semester = 'Spring'**  
**and year = 2009);**

## Join...using

- To provide the benefit of natural join while avoiding the danger of equating attributes erroneously, SQL provides a form of the natural join construct that allows you to specify exactly which columns should be equated. This feature is illustrated by the following query:
- ***select name, title from (instructor natural join teaches) join course using (course id);***
- The operation **join . . . using** requires a list of attribute names to be specified. Both inputs must have attributes with the specified names.
- Consider the operation ***r1 join r2 using(A1, A2).***

## Join...on

- The **on** condition allows a general predicate over the relations being joined. This predicate is written like a **where** clause predicate except for the use of the keyword **on** rather than **where**. Like the **using** condition, the **on** condition appears at the end of the join expression.
- Consider the following query, which has a join expression containing the **on** condition.
- **select \* from student join takes on student.ID= takes.ID;**
- The **on** condition above specifies that a tuple from *student* matches a tuple from *takes* if their *ID* values are equal.

- The following SQL statement will select all customers, and any orders they might have:
- ```
SELECT Customers.CustomerName, Orders.OrderID
FROM Customers
LEFT OUTER JOIN Orders ON Customers.CustomerID =
Orders.CustomerID;
```

# Inner join

- To distinguish normal joins from outer joins, normal joins are called **inner joins** in SQL. A join clause can thus specify **inner join** instead of **outer join** to specify that a normal join is to be used. The keyword **inner** is, however, optional.
- The default join type, when the **join** clause is used without the **outer** prefix is the **inner join**.

~~select \* from student join takes using (ID);~~

is equivalent to:

**select \* from student inner join takes using (ID);**

- The INNER JOIN keyword selects records that have matching values in both tables.
- SELECT column\_name(s)  
FROM table1  
INNER JOIN table2 using non attribute is considered if not  
ON table1.column\_name = table2.column\_name;

# Joined Relations – Examples

*select course\_id from*

- course **inner join** prereq **on**  
 $course.course\_id = prereq.course\_id$

*course*

| course_id | title       | dept_name  | credits | prereq_id | course_id |
|-----------|-------------|------------|---------|-----------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   | BIO-301   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    | CS-190    |

*prereq*

*→ Natural Join*

*Depends on the  $\exists \forall$*

*natural join → may result in  $\equiv$*

*→ shows*

- What is the difference between the above, and a natural join?  
\_\_\_\_\_
- Notice that we do not repeat those attributes that appear in the schemas of both relations; rather they appear only once.  
\_\_\_\_\_
- Notice also the order in which the attributes are listed: first the attributes common to the schemas of both relations, second those attributes unique to the schema of the first relation, and finally, those attributes unique to the schema of the second relation.

# Joined Relations – Examples

n course ~~natural right outer join~~ ~~prereq~~

| course_id | title       | dept_name  | credits | prereq_id |
|-----------|-------------|------------|---------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    |
| CS-347    | null        | null       | null    | CS-101    |

inner join  
vs

n

course **full outer join** prereq **using (course\_id)**

| course_id | title       | dept_name  | credits | prereq_id |
|-----------|-------------|------------|---------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    |
| CS-315    | Robotics    | Comp. Sci. | 3       | null      |
| CS-347    | null        | null       | null    | CS-101    |

{

# Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know an instructor's name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- A view provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a "virtual relation" is called a view.

# Does view occupy memory????

- Actually view is a stored select statement and its code is stored in data dictionary.
- It doesn't occupy space in the memory.
- whenever view is called, the stored select statement executes and fetches the data from the base table.
- If base table is deleted, view becomes invalid.

*Inception*  
*Update*

# View Definition

A view is defined using the **create view** statement which has the form

*name of the view*  
**create view  $v$  as <query expression>**

*select statement*

where <query expression> is any legal SQL expression. The view name is represented by  $v$ .

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
  - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view.

*No storage is created  
for view :- virtual*

# Example Views

A view of instructors without their salary

~~create view faculty as~~  
select ID, name, dept\_name  
from instructor

- Find all instructors in the Biology department

select name  
from faculty  
where dept\_name = 'Biology'

- The attribute names of a view can be specified explicitly as follows:

- To Create a view of department salary totals

~~create view departments\_total\_salary(dept\_name, total\_salary) as~~  
select dept\_name, sum (salary)  
from instructor  
group by dept\_name;

# Views Defined Using Other Views

n create view physics\_fall\_2009 as

```
select course.course_id, sec_id, building, room_number  
from course, section  
where course.course_id = section.course_id  
and course.dept_name = 'Physics'  
and section.semester = 'Fall'  
and section.year = 2009;
```

*list of courses  
by Physics dept  
in Fall 2009*

n create view physics\_fall\_2009\_watson as

```
select course_id, room_number  
from physics_fall_2009  
where building = 'Watson';
```

# View Expansion

- Expand use of a view in a query/another view
- ```
create view physics_fall_2009_watson as  
  (select course_id, room_number  
   from (select course.course_id, building, room_number  
         from course, section  
        where course.course_id = section.course_id  
          and course.dept_name = 'Physics'  
          and section.semester = 'Fall'  
          and section.year = '2009')  
    where building= 'Watson';
```
- (Note: The original image contains several red hand-drawn annotations. A large oval encloses the entire query. Red lines point from the word 'view' in the first line to the opening parenthesis of the subquery, and from the word 'another' in the list item to the closing parenthesis of the subquery. A bracket on the right side groups the subquery and the final WHERE clause, with the label 'another table.' written next to it.)*

## modification

# Updating a view

- The difficulty is that a modification to the database expressed in terms of a view must be translated to a modification to the actual relations in the logical model of the database.
- ~~create view faculty as base table~~  
~~select ID, name, dept\_name~~  
~~from instructor~~

*insert  
delete  
update*

- If a tuple is added as follows:  

- insert into faculty values ('30765', 'Green', 'Music');
- This insertion must be represented by an insertion into the relation *instructor*, since *instructor* is the actual relation from which the database system constructs the view *faculty*. However, to insert a tuple into *instructor*, we must have some value for *salary*.

- There are two reasonable approaches to dealing with this insertion:  
~~rejecting~~  
~~accepting~~
- Reject the insertion, and return an error message to the user.
- Insert a tuple ('30765', 'Green', 'Music', *null*) into the *instructor* relation.

~~A don't really~~

# Update of a View

- Add a new tuple to *faculty* view which we defined earlier

*user writes*  
`insert into faculty values ('30765', 'Green', 'Music');`

This insertion must be represented by the insertion of the tuple

*SQL writes*  
`insert into instructor values ('30765', 'Green',  
'Music', null)`  
into the *instructor* relation

## Some Updates cannot be Translated Uniquely

~~create view instructor\_info as~~

~~select ID, name, building  
from instructor, department  
where instructor.dept\_name= department.dept\_name;~~

- ~~insert into instructor\_info values ('69987', 'White', 'AB5');~~
  - If id does not exists
  - what if no department is in AB5?
- Suppose we Insert into ~~instructor~~ the ( 69987, 'White', null, null)
- ~~Insert in department(null, AB5, null)~~
- In the absence of common dept name this view update will fail



P. X.

FIX.

list relations

# Updating main table through views

A screenshot of the Oracle SQL Developer interface. The top window is a 'Worksheet' showing the SQL command 'select \* from dept;'. Below it is a 'Query Result' window displaying the following data:

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON

A screenshot of the Oracle SQL Developer interface. The top window is a 'Worksheet' showing the SQL command to create a view:

```
create view v_dept  
as  
select dname, loc from DEPT;
```

The bottom window is a 'Script Output' window showing the message 'view V\_DEPT created.'

Worksheet      Query Builder

```
select * from v_dept;
```

Script Output      Query Result

SQL | All Rows Fetched

	DNAME	LOC
1	ACCOUNTING	NEW YORK
2	RESEARCH	DALLAS
3	SALES	CHICAGO
4	OPERATIONS	BOSTON

Start Page      SCOTT

Worksheet      Query Builder

```
update v_dept  
set loc = 'USA'  
where dname = 'ACCOUNTING';
```

Script Output      Query Result

Task completed in 0.002 seconds

1 rows updated.

Worksheet      Query Builder

```
select * from v_dept;
```

Query Result

SQL | All Rows Fetched:

	DNAME	LOC
1	ACCOUNTING	USA
2	RESEARCH	DALLAS
3	SALES	CHICAGO
4	OPERATIONS	BOSTON

The screenshot shows the Oracle SQL Developer interface. The top menu bar includes icons for running, saving, and zooming. Below the menu is a toolbar with various buttons. The main window has tabs for 'Worksheet' and 'Query Builder', with 'Worksheet' selected. In the worksheet area, the following SQL command is entered:

```
select * from DEPT;
```

Below the worksheet, there are two tabs: 'Script Output' and 'Query Result'. The 'Query Result' tab is active, showing the output of the query. The output shows four rows of data from the DEPT table:

	DEPTNO	DNAME	LOC
1	10	ACCOUNTING	USA
2	20	RESEARCH	DALLAS
3	30	SALES	CHICAGO
4	40	OPERATIONS	BOSTON

Delete a view;

Drop viewname;

- 
- The diagram illustrates the creation of a view. At the top, there are two separate horizontal lines representing relations, labeled  $R_1$  and  $R_2$ . An arrow points from  $R_1$  to a circled area below it, which is labeled  $V$  (View). Another arrow points from  $R_2$  to the same circled area  $V$ .
- an SQL view is said to be **updatable** (that is, inserts, updates or deletes can be applied on the view) if the following conditions are all satisfied by the query defining the view:

- The **from** clause has only one database relation.
- The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or distinct specification.
- Any attribute not listed in the select clause can be set to null
- The query does not have a group by or having clause.

aggregate functions

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000
33456	Gold	Physics	87000
45565	Katz	Comp. Sci.	75000
58583	Califieri	History	62000
76543	Singh	Finance	80000
76766	Crick	Biology	72000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
69987	White	<i>null</i>	<i>null</i>

*instructor*

<i>dept_name</i>	<i>building</i>	<i>budget</i>
Biology	Watson	90000
Comp. Sci.	Taylor	100000
Elec. Eng.	Taylor	85000
Finance	Painter	120000
History	Painter	50000
Music	Packard	80000
Physics	Watson	70000
<i>null</i>	Taylor	<i>null</i>

*department*

# with check option

- ```
create view history_instructors as
    select *
      from instructor
     where dept_name= 'History';
```
- What happens if we insert ('25566', 'Brown', 'Biology', 100000) into *history\_instructors*?

- By default, SQL would allow the above update to proceed.
- However, views can be defined with a **with check option** clause at the end of the view definition; then, if a tuple inserted into the view does not satisfy the view's **where** clause condition, the insertion is rejected by the database system.
- Updates are similarly rejected if the new value does not satisfy the **where** clause conditions.

```
1 CREATE [REDACTED] VIEW vps AS
2     SELECT
3         employeeNumber,
4         lastName,
5         firstName,
6         jobTitle,
7         extension,
8         email,
9         officeCode,
10        reportsTo
11    FROM
12        employees
13    WHERE
14        jobTitle LIKE '%VP%'
15    WITH CHECK OPTION;
```

```
1 INSERT INTO vps(employeeNumber,firstname,lastname,jobtitle,extension,email,office  
Code,reportsTo)  
2 VALUES(1704,'John','Smith','IT Staff','x9112','johnsmith@classicmodelcars.com',1,  
1703);
```

This time, MySQL rejected the insert and issued the following error message:

```
1 Error Code: 1369. CHECK OPTION failed 'classicmodels.vps'
```

# Materialized Views

- **Materializing a view**: create a physical table containing all the tuples in the result of the query defining the view
- If relations used in the query are updated, the **materialized view result becomes out of date**
  - Need to **maintain** the view, by updating the view whenever the underlying relations are updated.